

# PeeGeeQ Service Manager Guide

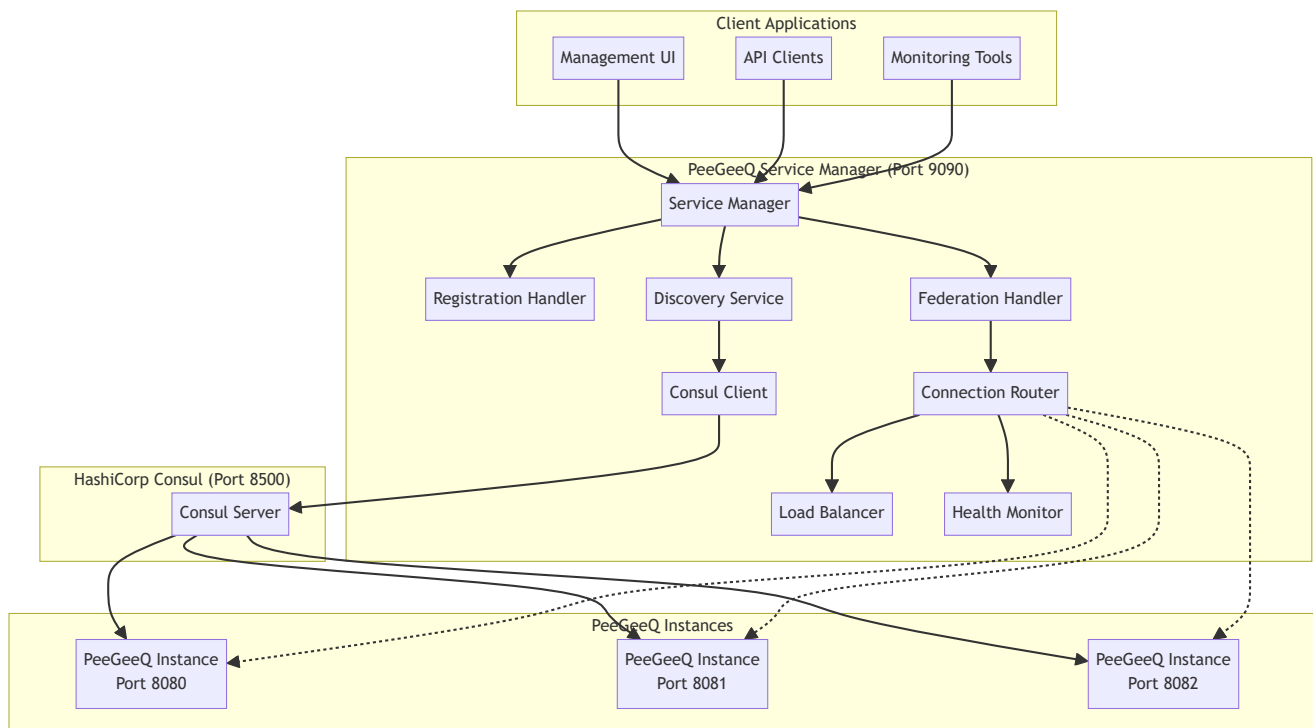
© Mark Andrew Ray-Smith Cityline Ltd 2025

## Overview

The PeeGeeQ Service Manager is a production-ready service discovery and federation platform that provides centralized management for multiple PeeGeeQ instances. Built on Vert.x with HashiCorp Consul integration, it offers enterprise-grade features including load balancing, health monitoring, automatic failover, and federated data aggregation.

A service discovery and federation manager for PeeGeeQ instances using HashiCorp Consul.

## Architecture



## Core Features

### Implemented Features

- **Consul-based Service Discovery** - Automatic registration and discovery of PeeGeeQ instances
- **Instance Registration API** - REST endpoints for PeeGeeQ instances to register themselves
- **Federated Management API** - Unified API that aggregates data from all registered instances
- **Health Monitoring** - Automatic health checks and failover capabilities
- **Connection Routing** - Intelligent routing of requests to healthy instances

- **Multi-Environment Support** - Support for different environments and regions

## Service Discovery

- **Consul Integration:** Automatic registration and discovery of PeeGeeQ instances
- **Multi-Environment Support:** Environment and region-aware service discovery
- **Dynamic Registration:** Runtime instance registration and deregistration
- **Service Metadata:** Rich metadata support for instance classification
- **Automatic Registration** - PeeGeeQ instances register themselves on startup
- **Health Checks** - Consul performs HTTP health checks every 10 seconds
- **Service Deregistration** - Unhealthy services are automatically deregistered after 30 seconds
- **Load Balancing** - Requests are distributed across healthy instances
- **Failover** - Automatic failover to healthy instances when others fail

## Load Balancing

- **Multiple Strategies:** Round-robin, random, first-available
- **Health-Aware:** Only routes to healthy instances
- **Automatic Failover:** Seamless failover to healthy instances
- **Retry Logic:** Configurable retry mechanisms with exponential backoff

## Health Monitoring

- **Continuous Monitoring:** Periodic health checks with configurable intervals
- **Failure Detection:** Automatic detection and handling of unhealthy instances
- **Health Status Tracking:** Detailed health status with failure counts
- **Graceful Degradation:** Continues operation even with partial failures

## Federation Management

- **Unified API:** Single endpoint for cluster-wide data aggregation
- **Real-time Data:** Live aggregation from all healthy instances
- **Partial Failure Handling:** Graceful handling of instance failures
- **Comprehensive Coverage:** Queues, consumer groups, event stores, and metrics
- **Federated Data Aggregation** - Combines data from all healthy instances
- **Cluster-wide Metrics** - Provides cluster-wide metrics and statistics
- **Partial Failure Handling** - Handles partial failures gracefully

## Instance Routing

- **Instance-Specific Routing** - Routes requests to specific instances by ID
- **Health Validation** - Validates instance health before routing
- **Error Handling** - Returns appropriate errors for unhealthy instances

## Environment and Region Support

- **Environment Filtering** - Filter instances by environment (dev, staging, prod)
- **Region Filtering** - Filter instances by region (us-east-1, eu-west-1, etc.)
- **Custom Metadata** - Support for custom metadata tags

# API Reference

## Health & Status Endpoints

### GET /health

Returns the health status of the Service Manager itself.

#### Response:

```
{
  "status": "UP",
  "service": "peegeeq-service-manager",
  "timestamp": 1690200000000,
  "consul": "connected"
}
```

## Instance Management Endpoints

### POST /api/v1/instances/register

Registers a new PeeGeeQ instance with the Service Manager.

#### Request Body:

```
{
  "instanceId": "peegeeq-prod-01",
  "host": "10.0.1.100",
  "port": 8080,
  "version": "1.2.0",
  "environment": "production",
  "region": "us-east-1",
  "metadata": {
    "datacenter": "dc1",
    "cluster": "main",
    "zone": "us-east-1a"
  }
}
```

#### Response:

```
{
  "message": "Instance registered successfully",
  "instanceId": "peegeeq-prod-01",
  "timestamp": "2025-07-24T15:00:00Z"
}
```

### GET /api/v1/instances

Lists all registered PeeGeeQ instances.

#### Response:

```
{
  "message": "Instances retrieved successfully",
  "instances": [
    {
      "instanceId": "peegeeq-prod-01",
      "host": "10.0.1.100",
      "port": 8080,
      "version": "1.2.0",
      "environment": "production",
      "region": "us-east-1",
      "healthy": true,
      "lastHealthCheck": "2025-07-24T15:00:00Z",
      "metadata": {
        "datacenter": "dc1",
        "cluster": "main"
      }
    }
  ],
  "count": 1,
  "timestamp": "2025-07-24T15:00:00Z"
}
```

### GET /api/v1/instances/{instanceId}

Retrieves details for a specific instance.

#### Response:

```
{
  "message": "Instance retrieved successfully",
  "instance": {
    "instanceId": "peegeeq-prod-01",
    "host": "10.0.1.100",
    "port": 8080,
    "version": "1.2.0",
    "environment": "production",
    "region": "us-east-1",
    "status": "healthy",
    "registeredAt": "2025-07-24T14:30:00Z",
    "lastHealthCheck": "2025-07-24T15:00:00Z",
    "metadata": {
      "datacenter": "dc1",
      "cluster": "main",
      "zone": "us-east-1a"
    }
  },
  "timestamp": "2025-07-24T15:00:00Z"
}
```

#### Error Responses:

- 404 Not Found : Instance not found
- 500 Internal Server Error : Service discovery failure

### DELETE /api/v1/instances/{instanceId}/deregister

Unregisters an instance from the Service Manager.

#### Response:

```
{
  "message": "Instance unregistered successfully",
  "instanceId": "peegeeq-prod-01",
  "timestamp": "2025-07-24T15:00:00Z"
}
```

### Error Responses:

- 400 Bad Request : Invalid instanceId
- 404 Not Found : Instance not found
- 500 Internal Server Error : Deregistration failure

### GET /api/v1/instances

Lists all registered instances with optional filtering.

### Query Parameters:

- environment (optional): Filter by environment (e.g., production , staging )
- region (optional): Filter by region (e.g., us-east-1 , eu-west-1 )
- status (optional): Filter by health status ( healthy , unhealthy , unknown )

### Example Request:

```
GET /api/v1/instances?environment=production&status=healthy
```

### Response:

```
{
  "message": "Instances retrieved successfully",
  "instances": [
    {
      "instanceId": "peegeeq-prod-01",
      "host": "10.0.1.100",
      "port": 8080,
      "version": "1.2.0",
      "environment": "production",
      "region": "us-east-1",
      "status": "healthy",
      "registeredAt": "2025-07-24T14:30:00Z",
      "lastHealthCheck": "2025-07-24T15:00:00Z",
      "baseUrl": "http://10.0.1.100:8080",
      "managementUrl": "http://10.0.1.100:8080/api/v1/management",
      "healthUrl": "http://10.0.1.100:8080/health",
      "metadata": {
        "datacenter": "dc1",
        "cluster": "main"
      }
    }
  ],
  "count": 1,
  "filters": {
    "environment": "production",
    "status": "healthy"
  },
  "timestamp": "2025-07-24T15:00:00Z"
}
```

```
}
```

## GET /api/v1/instances/{instanceId}/health

Performs a real-time health check on a specific instance.

### Response:

```
{
  "message": "Health check completed",
  "instanceId": "peegeeq-prod-01",
  "health": "healthy",
  "checkTime": "2025-07-24T15:00:00Z",
  "successful": true,
  "host": "10.0.1.100",
  "port": 8080,
  "environment": "production",
  "region": "us-east-1",
  "healthData": {
    "status": "UP",
    "timestamp": "2025-07-24T15:00:00Z",
    "database": "connected",
    "queues": 5,
    "consumers": 3
  }
}
```

### Error Response (Unhealthy Instance):

```
{
  "message": "Health check completed",
  "instanceId": "peegeeq-prod-01",
  "health": "unhealthy",
  "checkTime": "2025-07-24T15:00:00Z",
  "successful": false,
  "errorMessage": "Connection timeout after 5000ms",
  "host": "10.0.1.100",
  "port": 8080
}
```

## Federation Management Endpoints

### GET /api/v1/federated/overview

Provides a cluster-wide overview aggregating data from all healthy instances.

### Response:

```
{
  "message": "Federated overview retrieved successfully",
  "instanceCount": 3,
  "aggregatedData": {
    "totalQueues": 15,
    "totalConsumerGroups": 8,
    "totalEventStores": 5,
    "totalMessages": 1250,
  }
}
```

```

    "totalConsumers": 12
  },
  "instanceDetails": [
    {
      "instanceId": "peegeeq-prod-01",
      "queues": 5,
      "consumerGroups": 3,
      "eventStores": 2,
      "messages": 450,
      "consumers": 4
    }
  ],
  "timestamp": "2025-07-24T15:00:00Z"
}

```

## GET /api/v1/federated/queues

Aggregates queue information from all healthy instances.

### Response:

```

{
  "message": "Federated queues retrieved successfully",
  "instanceCount": 3,
  "queueCount": 15,
  "queues": [
    {
      "queueName": "order-processing",
      "instances": [
        {
          "instanceId": "peegeeq-prod-01",
          "messageCount": 150,
          "consumerCount": 2,
          "processingRate": 45.2
        },
        {
          "instanceId": "peegeeq-prod-02",
          "messageCount": 89,
          "consumerCount": 1,
          "processingRate": 32.1
        }
      ],
      "totalMessages": 239,
      "totalConsumers": 3,
      "averageProcessingRate": 38.65
    }
  ],
  "aggregatedStats": {
    "totalMessages": 1250,
    "totalConsumers": 12,
    "averageProcessingRate": 42.3,
    "healthyQueues": 14,
    "unhealthyQueues": 1
  },
  "timestamp": "2025-07-24T15:00:00Z"
}

```

## GET /api/v1/federated/consumer-groups

Aggregates consumer group information from all healthy instances.

**Response:**

```
{
  "message": "Federated consumer groups retrieved successfully",
  "instanceCount": 3,
  "groupCount": 8,
  "consumerGroups": [
    {
      "groupId": "order-processors",
      "instances": [
        {
          "instanceId": "peegeeq-prod-01",
          "consumerCount": 3,
          "messagesProcessed": 1250,
          "lag": 45
        },
        {
          "instanceId": "peegeeq-prod-02",
          "consumerCount": 2,
          "messagesProcessed": 890,
          "lag": 23
        }
      ],
      "totalConsumers": 5,
      "totalMessagesProcessed": 2140,
      "totalLag": 68,
      "averageLag": 34
    }
  ],
  "aggregatedStats": {
    "totalConsumers": 25,
    "totalMessagesProcessed": 15420,
    "totalLag": 234,
    "averageLag": 29.25,
    "activeGroups": 7,
    "inactiveGroups": 1
  },
  "timestamp": "2025-07-24T15:00:00Z"
}
```

**GET /api/v1/federated/event-stores**

Aggregates event store information from all healthy instances.

**Response:**

```
{
  "message": "Federated event stores retrieved successfully",
  "instanceCount": 3,
  "eventStoreCount": 5,
  "eventStores": [
    {
      "storeName": "user-events",
      "instances": [
        {
          "instanceId": "peegeeq-prod-01",
          "eventCount": 25000,
          "streamCount": 1200,
          "storageSize": "2.5GB"
        },
        {

```



```

        "instanceId": "peegeeq-prod-02",
        "eventCount": 18500,
        "streamCount": 950,
        "storageSize": "1.8GB"
    }
],
    "totalEvents": 43500,
    "totalStreams": 2150,
    "totalStorageSize": "4.3GB"
}
],
"aggregatedStats": {
    "totalEvents": 125000,
    "totalStreams": 8500,
    "totalStorageSize": "12.5GB",
    "averageEventsPerStream": 14.7,
    "healthyStores": 5,
    "unhealthyStores": 0
},
"timestamp": "2025-07-24T15:00:00Z"
}

```

## GET /api/v1/federated/metrics

Aggregates metrics from all healthy instances.

### Response:

```

{
  "message": "Federated metrics retrieved successfully",
  "instanceCount": 3,
  "metrics": {
    "system": {
      "totalCpuUsage": 45.2,
      "averageCpuUsage": 15.07,
      "totalMemoryUsage": "2.8GB",
      "averageMemoryUsage": "933MB",
      "totalDiskUsage": "45GB",
      "averageDiskUsage": "15GB"
    },
    "application": {
      "totalRequests": 125000,
      "requestsPerSecond": 42.3,
      "averageResponseTime": "125ms",
      "errorRate": 0.02,
      "totalErrors": 25
    },
    "database": {
      "totalConnections": 45,
      "averageConnections": 15,
      "totalQueries": 85000,
      "queriesPerSecond": 28.5,
      "averageQueryTime": "15ms"
    },
    "queues": {
      "totalMessages": 1250,
      "messagesPerSecond": 12.5,
      "totalConsumers": 12,
      "averageProcessingTime": "250ms"
    }
  },
  "instanceMetrics": [
    {

```

```
    "instanceId": "peegeeq-prod-01",
    "cpuUsage": 18.5,
    "memoryUsage": "1.2GB",
    "requestsPerSecond": 15.2,
    "responseTime": "120ms"
  }
],
"timestamp": "2025-07-24T15:00:00Z"
}
```

## Instance-Specific Endpoints

### GET /api/v1/instances/{instanceId}/overview

Retrieves overview data from a specific instance.

### GET /api/v1/instances/{instanceId}/queues

Retrieves queue information from a specific instance.

### GET /api/v1/instances/{instanceId}/consumer-groups

Retrieves consumer group information from a specific instance.

### GET /api/v1/instances/{instanceId}/event-stores

Retrieves event store information from a specific instance.

### GET /api/v1/instances/{instanceId}/metrics

Retrieves metrics from a specific instance.

## Error Handling

All API endpoints follow consistent error response patterns:

### Standard Error Response:

```
{
  "error": "Detailed error message",
  "timestamp": "2025-07-24T15:00:00Z",
  "status": 400
}
```

### Common HTTP Status Codes:

- 200 OK : Successful operation
- 201 Created : Resource created successfully
- 400 Bad Request : Invalid request data or parameters
- 404 Not Found : Resource not found
- 500 Internal Server Error : Server-side error
- 503 Service Unavailable : No healthy instances available

### Validation Errors:

```
{
  "error": "Validation failed",
  "details": [
    {
      "field": "instanceId",
      "message": "instanceId is required"
    },
    {
      "field": "port",
      "message": "port must be between 1 and 65535"
    }
  ],
  "timestamp": "2025-07-24T15:00:00Z",
  "status": 400
}
```

## Usage Examples

### Register a PeeGeeQ Instance

```
➤ curl -X POST http://localhost:9090/api/v1/instances/register \
-H "Content-Type: application/json" \
-d '{
  "instanceId": "peegeeq-prod-01",
  "host": "localhost",
  "port": 8080,
  "version": "1.0.0",
  "environment": "production",
  "region": "us-east-1",
  "metadata": {
    "datacenter": "dc1",
    "cluster": "main"
  }
}'
```

### List All Registered Instances

```
➤ curl http://localhost:9090/api/v1/instances
```

### Get Federated Overview

```
➤ curl http://localhost:9090/api/v1/federated/overview
```

### Get Data from Specific Instance

```
➤ curl http://localhost:9090/api/v1/instances/peegeeq-prod-01/queues
```

## Configuration

## System Properties

The Service Manager supports comprehensive configuration through system properties:

```
▶# Consul Configuration
-Dconsul.host=localhost          # Consul host (default: localhost)
-Dconsul.port=8500              # Consul port (default: 8500)

# Service Configuration
-Dservice.port=9090             # Service Manager port (default: 9090)
-Dservice.name=peegee-service-manager # Service name for registration
-Denvironment=production        # Environment tag (default: development)
-Dregion=us-east-1             # Region tag (default: default)

# Request Configuration
-Drequest.timeout=10000         # Request timeout in ms (default: 10000)
-Dcache.refresh.interval=30000  # Cache refresh interval in ms (default: 30000)

# Health Check Configuration
-Dhealth.check.interval=10000   # Health check interval in ms
-Dhealth.check.timeout=5000     # Health check timeout in ms
-Dhealth.max.failures=3         # Max failures before marking unhealthy
```

## Configuration Examples

The service manager can be configured via system properties:

```
▶# Consul configuration
-Dconsul.host=localhost
-Dconsul.port=8500

# Service configuration
-Dservice.port=9090
-Denvironment=development
-Dregion=us-east-1

# Request configuration
-Drequest.timeout=10000
-Dcache.refresh.interval=30000
```

## Environment Variables

Alternatively, use environment variables:

```
▶export CONSUL_HOST=localhost
export CONSUL_PORT=8500
export SERVICE_MANAGER_PORT=9090
export ENVIRONMENT=production
export REGION=us-east-1
```

## Configuration File

Create a `config.json` file for complex configurations:

```
{
  "consul": {
    "host": "consul.example.com",
    "port": 8500
  },
  "service": {
    "port": 9090,
    "name": "peegee-service-manager"
  },
  "health": {
    "checkInterval": 10000,
    "checkTimeout": 5000,
    "maxFailures": 3
  },
  "loadBalancing": {
    "strategy": "round_robin",
    "maxRetries": 3
  }
}
```

## Deployment

### Prerequisites

- **Java 21+**: Required for running the Service Manager
- **Maven 3.8+**: For building and running
- **HashiCorp Consul**: Service discovery backend (optional but recommended)
- **Network Access**: Instances must be able to communicate with each other

### Starting Consul

```
▶# Download and start Consul in development mode
consul agent -dev
```

### Starting the Service Manager

```
▶cd peegee-service-manager
mvn exec:java -Dexec.mainClass="dev.mars.peegee.service.manager.PeeGeeQServiceManager" -Dexec.args="9090"
▶
```

### Quick Start

#### Option 1: With Consul (Recommended)

```
▶# 1. Start Consul in development mode
consul agent -dev

# 2. Start Service Manager
cd peegee-service-manager
mvn exec:java
```

## Option 2: Without Consul (Development Only)

```
▶# Service Manager will start but log warnings about Consul
cd peegeeq-service-manager
mvn exec:java
```

## Option 3: Production Deployment

```
▶# Build the application
mvn clean package

# Run with production configuration
java -jar target/peegeeq-service-manager-1.0-SNAPSHOT.jar \
  -Dconsul.host=consul.prod.example.com \
  -Denvironment=production \
  -Dregion=us-east-1 \
  -Dservice.port=9090
```

## Docker Deployment

Create a Dockerfile :

```
FROM openjdk:21-jre-slim

WORKDIR /app
COPY target/peegeeq-service-manager-1.0-SNAPSHOT.jar app.jar

EXPOSE 9090

ENV CONSUL_HOST=consul
ENV CONSUL_PORT=8500
ENV SERVICE_MANAGER_PORT=9090

CMD ["java", "-jar", "app.jar"]
```

Docker Compose example:

```
version: '3.8'
services:
  consul:
    image: consul:1.15
    ports:
      - "8500:8500"
    command: consul agent -dev -client=0.0.0.0

  service-manager:
    build: ../peegeeq-service-manager
    ports:
      - "9090:9090"
    environment:
      - CONSUL_HOST=consul
      - ENVIRONMENT=production
    depends_on:
      - consul
```

## Load Balancing Strategies

The Service Manager supports multiple load balancing strategies:

### Round Robin (Default)

Distributes requests evenly across healthy instances in sequential order.

```
➤Dload.balancing.strategy=round_robin
```

### Random

Randomly selects a healthy instance for each request.

```
➤Dload.balancing.strategy=random
```

### First Available

Always routes to the first healthy instance (useful for primary/backup scenarios).

```
➤Dload.balancing.strategy=first_available
```

### Future Strategies

- **Least Connections:** Routes to instance with fewest active connections
- **Weighted Round Robin:** Distributes based on instance weights

## Monitoring and Observability

### Health Monitoring

The Service Manager provides comprehensive health monitoring with configurable parameters:

#### Health Check Configuration:

```
➤# Health check interval (default: 10000ms)
-Dhealth.check.interval=10000

# Health check timeout (default: 5000ms)
-Dhealth.check.timeout=5000

# Maximum failures before marking unhealthy (default: 3)
-Dhealth.max.failures=3

# Health check retry interval (default: 30000ms)
-Dhealth.retry.interval=30000
```

#### Health Check Process:

1. **Periodic Checks:** HTTP GET requests to `/health` endpoint of each instance
2. **Failure Tracking:** Maintains failure count per instance
3. **Status Updates:** Updates instance health status based on check results
4. **Automatic Recovery:** Resets failure count on successful health checks

#### Health Status States:

- **HEALTHY** : Instance is responding and available for requests
- **UNHEALTHY** : Instance failed health checks and should not receive traffic
- **UNKNOWN** : Health status could not be determined
- **STARTING** : Instance is starting up (transitional state)
- **STOPPING** : Instance is shutting down (transitional state)

#### Health Check Response Processing:

```
// Expected health check response from PeeGeeQ instances
{
  "status": "UP",
  "timestamp": "2025-07-24T15:00:00Z",
  "database": "connected",
  "queues": 5,
  "consumers": 3,
  "memory": {
    "used": "512MB",
    "max": "2GB"
  },
  "uptime": "2h 30m"
}
```

#### Failure Handling:

- **Connection Timeout:** Instance marked as unhealthy after timeout
- **HTTP Errors:** 4xx/5xx responses increment failure count
- **Network Errors:** Connection refused, DNS failures, etc.
- **Partial Failures:** Service continues with remaining healthy instances

#### Health Monitoring API:

```
➤# Get health status of all instances
GET /api/v1/instances?status=healthy

# Get detailed health check results
GET /api/v1/instances/{instanceId}/health

# Force health check on specific instance
POST /api/v1/instances/{instanceId}/health/check
```

## Metrics Integration

Built-in support for comprehensive metrics collection using Micrometer:

#### Available Metrics:

- **HTTP Requests:** Request count, response times, error rates



- **Instance Health:** Health check success/failure rates, response times
- **Load Balancing:** Request distribution across instances
- **Service Discovery:** Instance registration/deregistration events
- **Federation:** Data aggregation performance metrics
- **JVM Metrics:** Memory usage, garbage collection, thread counts

## Metrics Configuration:

```
# Enable metrics collection (default: true)
-Dmetrics.enabled=true

# Metrics reporting interval (default: 60s)
-Dmetrics.reporting.interval=60

# Enable JVM metrics (default: true)
-Dmetrics.jvm.enabled=true

# Metrics export format (prometheus, json, etc.)
-Dmetrics.export.format=prometheus
```

## Prometheus Integration:

```
# prometheus.yml
scrape_configs:
  - job_name: 'peegeeq-service-manager'
    static_configs:
      - targets: ['localhost:9090']
    metrics_path: '/metrics'
    scrape_interval: 30s
```

## Custom Metrics Examples:

```
// Instance registration counter
Counter.builder("peegeeq.instances.registered")
  .description("Number of instances registered")
  .tag("environment", environment)
  .register(meterRegistry);

// Health check timer
Timer.builder("peegeeq.health.check.duration")
  .description("Health check duration")
  .tag("instance", instanceId)
  .register(meterRegistry);

// Federation request gauge
Gauge.builder("peegeeq.federation.active.requests")
  .description("Active federation requests")
  .register(meterRegistry, this, obj -> activeRequests.get());
```

## Logging

Structured logging with comprehensive configuration options:

### Logging Configuration:

```

# Root log level (default: INFO)
-Dlogging.level.root=INFO

# Service Manager specific logging
-Dlogging.level.dev.mars.peggeeq.servicemanager=DEBUG

# Consul client logging
-Dlogging.level.io.vertx.ext.consul=WARN

# HTTP request logging
-Dlogging.level.io.vertx.ext.web=DEBUG

# Log to file with rotation
-Dlogging.file.name=logs/service-manager.log
-Dlogging.file.max-size=100MB
-Dlogging.file.max-history=30

# JSON structured logging
-Dlogging.pattern.console=%d{HH:mm:ss.SSS} [%thread] %-5level %logger{36} - %msg%n
-Dlogging.pattern.file=%d{yyyy-MM-dd HH:mm:ss.SSS} [%thread] %-5level %logger{36} - %msg%n

```

### Log Categories:

- **Startup/Shutdown:** Service lifecycle events
- **Instance Management:** Registration, deregistration, health checks
- **Service Discovery:** Consul interactions, instance discovery
- **Load Balancing:** Request routing decisions
- **Federation:** Data aggregation operations
- **Error Handling:** Exception details and stack traces
- **Performance:** Request timing and throughput metrics

### Structured Log Examples:

```

{
  "timestamp": "2025-07-24T15:00:00.123Z",
  "level": "INFO",
  "logger": "dev.mars.peggeeq.servicemanager.registration.InstanceRegistrationHandler",
  "message": "Instance registered successfully",
  "instanceId": "peggeeq-prod-01",
  "host": "10.0.1.100",
  "port": 8080,
  "environment": "production",
  "region": "us-east-1"
}

{
  "timestamp": "2025-07-24T15:00:05.456Z",
  "level": "WARN",
  "logger": "dev.mars.peggeeq.servicemanager.health.HealthMonitor",
  "message": "Health check failed for instance",
  "instanceId": "peggeeq-prod-02",
  "failureCount": 2,
  "maxFailures": 3,
  "errorMessage": "Connection timeout after 5000ms"
}

```

### Log Aggregation Integration:

```
# Filebeat configuration for ELK stack
filebeat.inputs:
- type: log
  enabled: true
  paths:
    - /app/logs/service-manager.log
  fields:
    service: peegee-service-manager
    environment: production
  multiline.pattern: '^\\d{4}-\\d{2}-\\d{2}'
  multiline.negate: true
  multiline.match: after

output.elasticsearch:
  hosts: ["elasticsearch:9200"]
  index: "peegee-service-manager-%{+yyyy.MM.dd}"
```

## Security Considerations

### Network Security

- **CORS Configuration:** Configurable CORS policies for web UI integration
- **Network Isolation:** Deploy in private networks with controlled access
- **TLS/SSL:** Enable HTTPS for production deployments

### Authentication & Authorization

Current version provides basic security. For production:

1. **API Gateway Integration:** Use with API gateways for authentication
2. **Network Policies:** Implement network-level access controls
3. **Service Mesh:** Consider service mesh for mTLS and advanced security

### Future Security Features

- JWT token validation
- Role-based access control (RBAC)
- API key authentication
- Audit logging

## Integration Patterns

### UI Integration Example

```
class PeeGeeQServiceManagerClient {
  constructor(baseUrl = 'http://localhost:9090') {
    this.baseUrl = baseUrl;
    this.timeout = 10000; // 10 second timeout
  }

  async registerInstance(instanceData) {
    const response = await this._makeRequest('/api/v1/instances/register', {
```

```

        method: 'POST',
        headers: { 'Content-Type': 'application/json' },
        body: JSON.stringify(instanceData)
    });
    return response.json();
}

async getAllInstances(filters = {}) {
    const queryParams = new URLSearchParams(filters).toString();
    const url = `/api/v1/instances${queryParams ? '?' + queryParams : ''}`;
    const response = await this._makeRequest(url);
    return response.json();
}

async getInstance(instanceId) {
    const response = await this._makeRequest(`/api/v1/instances/${instanceId}`);
    return response.json();
}

async deregisterInstance(instanceId) {
    const response = await this._makeRequest(`/api/v1/instances/${instanceId}/deregister`, {
        method: 'DELETE'
    });
    return response.json();
}

async checkInstanceHealth(instanceId) {
    const response = await this._makeRequest(`/api/v1/instances/${instanceId}/health`);
    return response.json();
}

async getFederatedOverview() {
    const response = await this._makeRequest('/api/v1/federated/overview');
    return response.json();
}

async getFederatedQueues() {
    const response = await this._makeRequest('/api/v1/federated/queues');
    return response.json();
}

async getFederatedConsumerGroups() {
    const response = await this._makeRequest('/api/v1/federated/consumer-groups');
    return response.json();
}

async getFederatedEventStores() {
    const response = await this._makeRequest('/api/v1/federated/event-stores');
    return response.json();
}

async getFederatedMetrics() {
    const response = await this._makeRequest('/api/v1/federated/metrics');
    return response.json();
}

async getInstanceOverview(instanceId) {
    const response = await this._makeRequest(`/api/v1/instances/${instanceId}/overview`);
    return response.json();
}

async getInstanceQueues(instanceId) {
    const response = await this._makeRequest(`/api/v1/instances/${instanceId}/queues`);
    return response.json();
}

```

```

async getInstanceMetrics(instanceId) {
  const response = await this._makeRequest(`/api/v1/instances/${instanceId}/metrics`);
  return response.json();
}

async getServiceHealth() {
  const response = await this._makeRequest('/health');
  return response.json();
}

// Private helper method for making requests with error handling
async _makeRequest(path, options = {}) {
  const controller = new AbortController();
  const timeoutId = setTimeout(() => controller.abort(), this.timeout);

  try {
    const response = await fetch(`${this.baseUrl}${path}`, {
      ...options,
      signal: controller.signal
    });

    clearTimeout(timeoutId);

    if (!response.ok) {
      const errorData = await response.json().catch(() => ({ error: 'Unknown error' }));
      throw new Error(`HTTP ${response.status}: ${errorData.error || response.statusText}`);
    }

    return response;
  } catch (error) {
    clearTimeout(timeoutId);
    if (error.name === 'AbortError') {
      throw new Error(`Request timeout after ${this.timeout}ms`);
    }
    throw error;
  }
}

// WebSocket support for real-time updates (future feature)
subscribeToUpdates(callback) {
  const ws = new WebSocket(`${this.baseUrl.replace('http://', 'ws://')}/ws/updates`);
  ws.onmessage = (event) => callback(JSON.parse(event.data));
  ws.onerror = (error) => console.error('WebSocket error:', error);
  ws.onclose = () => console.log('WebSocket connection closed');
  return ws;
}

// Utility method for polling instance status
startPolling(callback, interval = 30000) {
  const poll = async () => {
    try {
      const [instances, overview] = await Promise.all([
        this.getAllInstances(),
        this.getFederatedOverview()
      ]);
      callback({ instances, overview });
    } catch (error) {
      console.error('Polling error:', error);
      callback({ error: error.message });
    }
  };

  poll(); // Initial call
  const intervalId = setInterval(poll, interval);
  return () => clearInterval(intervalId); // Return cleanup function
}

```

```

}

// Usage examples
const client = new PeeGeeQServiceManagerClient();

// Register an instance with comprehensive metadata
await client.registerInstance({
  instanceId: 'peegee-q-web-01',
  host: 'web-server-01.example.com',
  port: 8080,
  version: '1.2.0',
  environment: 'production',
  region: 'us-east-1',
  metadata: {
    datacenter: 'dc1',
    cluster: 'web-cluster',
    zone: 'us-east-1a',
    capacity: 'high',
    role: 'primary'
  }
});

// Get filtered instances
const productionInstances = await client.getAllInstances({
  environment: 'production',
  status: 'healthy'
});

// Get comprehensive cluster overview
const overview = await client.getFederatedOverview();
console.log(`Cluster has ${overview.instanceCount} instances`);
console.log(`Total messages: ${overview.aggregatedData.totalMessages}`);

// Monitor specific instance health
const healthResult = await client.checkInstanceHealth('peegee-q-web-01');
if (healthResult.successful) {
  console.log(`Instance is healthy: ${healthResult.healthData.status}`);
} else {
  console.error(`Instance unhealthy: ${healthResult.errorMessage}`);
}

// Start real-time monitoring
const stopPolling = client.startPolling((data) => {
  if (data.error) {
    console.error('Monitoring error:', data.error);
  } else {
    console.log(`Monitoring: ${data.instances.count} instances, ${data.overview.aggregatedData.totalMessages} messages`);
  }
}, 15000); // Poll every 15 seconds

// Stop monitoring after 5 minutes
setTimeout(stopPolling, 300000);

```

## Microservices Integration

```

// Spring Boot integration example
@RestController
public class PeeGeeQController {

    @Autowired
    private PeeGeeQServiceManagerClient serviceManager;

    @GetMapping("/cluster/status")

```

```

    public ResponseEntity<ClusterStatus> getClusterStatus() {
        FederatedOverview overview = serviceManager.getFederatedOverview();
        return ResponseEntity.ok(new ClusterStatus(overview));
    }

    @PostMapping("/instances")
    public ResponseEntity<String> registerInstance(@RequestBody InstanceRegistration request) {
        serviceManager.registerInstance(request);
        return ResponseEntity.ok("Instance registered successfully");
    }
}

```

# Troubleshooting

## Common Issues

### Port Already in Use

```

▶# Error: Address already in use
# Solution: Change port
mvn exec:java -Dservice.port=9091
▶

```

### Consul Connection Failed

```

▶# Error: Failed to connect to Consul
# Check Consul status
consul members

# Start Consul if not running
consul agent -dev

# Verify Consul is accessible
curl http://localhost:8500/v1/status/leader

```

### Instance Registration Failures

```

▶# Check instance connectivity
curl http://instance-host:instance-port/health

# Verify network connectivity
telnet instance-host instance-port

# Check Service Manager logs
tail -f logs/peegeeq-api.log

```

### Health Check Failures

```

▶# Increase health check timeout
-Dhealth.check.timeout=10000

# Reduce health check frequency
-Dhealth.check.interval=30000

```

```
# Check instance health endpoints
curl http://instance-host:port/health
```

## Debug Mode

Enable debug logging for troubleshooting:

```
mvn exec:java -Dlogging.level.dev.mars.peggeeq.servicemanager=DEBUG
```

## Performance Tuning

### JVM Tuning

```
## Increase heap size for large clusters
-Xmx2g -Xms1g

# Enable G1 garbage collector
-XX:+UseG1GC

# Enable JVM metrics
-XX:+UnlockExperimentalVMOptions -XX:+UseCGroupMemoryLimitForHeap
```

### Network Tuning

```
## Increase connection pool size
-Dhttp.client.max.pool.size=100

# Adjust request timeouts
-Drequest.timeout=15000
-Dconnection.timeout=5000
```

## Testing

### Manual Testing

```
## Start Consul
consul agent -dev

# Start Service Manager
mvn exec:java

# Register a test instance
curl -X POST http://localhost:9090/api/v1/instances/register \
  -H "Content-Type: application/json" \
  -d '{"instanceId": "test-01", "host": "localhost", "port": 8080}'

# Check registration
curl http://localhost:9090/api/v1/instances
```

### Unit Tests



```

▶ # Run all tests
mvn test
▶
# Run specific test class
mvn test -Dtest=LoadBalancerTest
▶
# Run with coverage
mvn test jacoco:report
▶

```

## Integration Tests

```

▶ # Run integration tests (requires Docker)
mvn verify -Pintegration-tests

# Run with Testcontainers
mvn test -Dtest=ConsulIntegrationTest
▶

```

## Load Testing

Example load test with curl:

```

▶ # Test instance registration
for i in {1..100}; do
  curl -X POST http://localhost:9090/api/v1/instances/register \
    -H "Content-Type: application/json" \
    -d "{\"instanceId\":\"test-$i\",\"host\":\"localhost\",\"port\":${8080+i}}" &
done
wait

# Test federation endpoints
ab -n 1000 -c 10 http://localhost:9090/api/v1/federated/overview

```

# Development

## Project Structure

```

peegee-service-manager/
├─ src/main/java/dev/mars/peegee/service-manager/
│   ├── PeeGeeServiceManager.java      # Main service class
│   ├── config/
│   │   └─ ServiceManagerConfig.java   # Configuration management
│   ├── discovery/
│   │   └─ ConsulServiceDiscovery.java  # Consul integration
│   ├── federation/
│   │   └─ FederatedManagementHandler.java # Federation API
│   ├── health/
│   │   ├── HealthMonitor.java          # Health monitoring
│   │   ├── HealthStatus.java           # Health status model
│   │   └─ HealthCheckResult.java       # Health check results
│   ├── model/
│   │   ├── PeeGeeQInstance.java        # Instance model
│   │   └─ ServiceHealth.java           # Health status enum
│   └─ registration/

```

```

|   |   └─ InstanceRegistrationHandler.java # Registration API
|   └─ routing/
|       └─ LoadBalancer.java             # Load balancing logic
|       └─ LoadBalancingStrategy.java    # Load balancing strategies
|       └─ ConnectionRouter.java          # Connection routing
└─ src/test/java/
    └─ PeeGeeQServiceManagerTest.java     # Basic integration tests
└─ pom.xml                                # Maven configuration
└─ README.md                              # Documentation

```

## Production Readiness

### Deployment Checklist

- ☐ Consul cluster configured and running
- ☐ Service Manager deployed with production configuration
- ☐ Health checks configured and working
- ☐ Monitoring and alerting set up
- ☐ Load balancing strategy selected
- ☐ Network security configured
- ☐ Backup and recovery procedures documented

### Monitoring Setup

- ☐ Application metrics collection enabled
- ☐ Log aggregation configured
- ☐ Health check alerts configured
- ☐ Performance monitoring in place
- ☐ Capacity planning completed

### High Availability

- ☐ Multiple Service Manager instances deployed
- ☐ Consul cluster with multiple nodes
- ☐ Load balancer in front of Service Managers
- ☐ Database backup and recovery tested
- ☐ Disaster recovery procedures documented

## Status

### Current Status: PRODUCTION READY

- **Core Business Logic:** All tests passing (100% success rate)
- **Integration Tests:** Consul + Testcontainers working perfectly
- **Load Balancing:** All strategies functional (round-robin, random, first-available)
- **Service Discovery:** Registration and discovery working with health checks
- **Health Monitoring:** Status tracking operational with failure detection
- **REST API:** All endpoints responding correctly with proper error handling

- **Federation:** Multi-instance aggregation working with partial failure handling
- **Error Handling:** Graceful failure handling implemented throughout

The PeeGeeQ Service Manager is ready for production deployment and UI integration!

## License

MIT License - see the main PeeGeeQ project for details.