

# PeeGeeQ Database Setup Guide

---

## Comprehensive guide to database initialization across all services and environments

Version: 1.0

Date: December 1, 2025

Author: Mark Andrew Ray-Smith Cityline Ltd

---

## Table of Contents

- [Overview](#)
  - [Architecture Summary](#)
  - [Schema Definition Sources and Alignment](#)
  - [Why Templates? Understanding the Design](#)
  - [How Templates Work](#)
  - [Production Setup](#)
  - [Development Setup](#)
  - [Test Setup](#)
  - [Service-Specific Patterns](#)
  - [Database Schema Components](#)
  - [Setup Methods Comparison](#)
  - [Troubleshooting](#)
- 

## Overview

PeeGeeQ uses different database initialization strategies depending on the environment and service:

Environment	Method	Tool	Description
<b>Production</b>	Flyway Migrations	<code>peegueq-migrations</code> module	Standalone JAR in CI/CD pipeline
<b>Development</b>	Flyway Migrations or Docker	Maven plugin or docker-compose	Automated or manual via scripts
<b>Integration Tests</b>	Programmatic Setup	<code>PeeGeeQDatabaseSetupService</code>	Template-based dynamic setup
<b>Unit Tests</b>	SQL Script Injection	<code>PeeGeeQTestSchemaInitializer</code>	Direct JDBC SQL execution

**Key Principle:** Production uses **migrations** (Flyway), tests use **programmatic setup** (templates).

**Important:** PeeGeeQ supports **custom PostgreSQL schemas** – the entire system can be deployed to any schema (e.g., `peegueq`, `myapp_queue`, `tenant_a`) instead of the default `public` schema. This enables multi-tenant deployments, namespace isolation, and integration with existing databases.

---

# Custom Schema Support

## Overview

PeeGeeQ is designed to work with **any PostgreSQL schema**, not just the default `public` schema. This enables:

- **Multi-tenant deployments:** `tenant_a`, `tenant_b`, `tenant_c` schemas
- **Namespace isolation:** `myapp_queue`, `yourapp_queue` in same database
- **Integration with existing systems:** Deploy PeeGeeQ into existing database schemas
- **Security segregation:** Different schemas with different permissions
- **Schema-based sharding:** Horizontal scaling using schema partitioning

## Default Schema Configuration

By default, PeeGeeQ uses **two schemas**:

Schema	Purpose	Contains
<code>peegeeq</code>	Main queue operations	<code>queue_messages</code> , <code>outbox</code> , <code>outbox_consumer_groups</code> , <code>dead_letter_queue</code> , <code>queue_template</code> , consumer group tables
<code>bitemporal</code>	Event sourcing	<code>bitemporal_event_log</code> , <code>event_store_template</code> , temporal views

## Why Two Schemas?

1. **Logical Separation:** Queue operations vs event sourcing are distinct concerns
2. **Permission Management:** Grant queue access without event store access
3. **Backup Strategies:** Backup event store separately from queues
4. **Schema Evolution:** Evolve queue schema independently from event schema

## Configuring Custom Schemas

### Production (Flyway Migrations)

#### Method 1: Flyway Configuration File

```
# peegeeq-migrations/src/main/resources/application.properties

# Example 1: Custom schema names
flyway.schemas=myapp_message,myapp_events

# Example 2: Single schema for everything
flyway.schemas=myapp

# Example 3: Multi-tenant with tenant ID
flyway.schemas=tenant_${TENANT_ID}_queue,tenant_${TENANT_ID}_events

# Example 4: Namespace isolation
flyway.schemas=acme_peegeeq,acme_bitemporal
```

## Method 2: Environment Variables

```
# Set custom schemas via environment
export FLYWAY_SCHEMAS="myapp.messaging,myapp_events"

# Run migrations
java -jar peegeeq-migrations.jar migrate
```

## Method 3: Command Line

```
# Maven plugin
mvn flyway:migrate \
-Dflyway.schemas=myapp.messaging,myapp_events

# Standalone JAR
java -jar peegeeq-migrations.jar migrate \
-schemas=myapp.messaging,myapp_events
```

## Application Configuration

### Java Configuration

```
// Option 1: Using PeeGeeQConfiguration constructor
PeeGeeQConfiguration config = new PeeGeeQConfiguration(
    "production",           // profile
    "localhost",            // host
    5432,                  // port
    "myapp_db",             // database
    "dbuser",               // username
    "password",             // password
    "myapp.messaging"       // CUSTOM SCHEMA
);

PeeGeeQManager manager = new PeeGeeQManager(config);
```

### Properties File

```
# application.properties
peegeeq.db.schema=myapp.messaging
peegeeq.db.host=localhost
peegeeq.db.port=5432
peegeeq.db.database=myapp_db
peegeeq.db.username=dbuser
peegeeq.db.password=password
```

## Environment Variables

```
export PEEGEEQ_DB_SCHEMA="myapp.messaging"
export PEEGEEQ_DB_HOST="localhost"
export PEEGEEQ_DB_PORT="5432"
export PEEGEEQ_DB_DATABASE="myapp_db"
export PEEGEEQ_DB_USERNAME="dbuser"
export PEEGEEQ_DB_PASSWORD="password"
```

## Integration Tests

```
@BeforeAll
void setupDatabase() throws Exception {
    setupService = new PeeGeeQDatabaseSetupService();

    // Custom schema in test
    DatabaseConfig dbConfig = new DatabaseConfig.Builder()
        .host(postgres.getHost())
        .port(postgres.getFirstMappedPort())
        .databaseName("test_db_" + UUID.randomUUID())
        .username(postgres.getUsername())
        .password(postgres.getPassword())
        .schema("myapp_test") // CUSTOM SCHEMA
        .build();

    DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
        .setupId("test_" + UUID.randomUUID())
        .databaseConfig(dbConfig)
        .queues(queueConfigs)
        .eventStores(eventStoreConfigs)
        .build();

    setupResult = setupService.createCompleteSetup(request).get();
}
```

## Schema Layout Examples

### Example 1: Single Schema for Everything

```
-- Use "myapp" for all PeeGeeQ tables
CREATE SCHEMA myapp;

-- All tables in myapp schema
myapp.queue_messages
myapp.outbox
myapp.outbox_consumer_groups
```

```
myapp.bitemporal_event_log -- event store also in myapp
myapp.dead_letter_queue
myapp.queue_template
myapp.event_store_template
-- ... all other tables
```

**Configuration:**

```
flyway.schemas=myapp
peegeq.db.schema=myapp
```

**Example 2: Separate Schemas for Queue and Events**

```
-- Queue operations in one schema
CREATE SCHEMA acme_queue;
acme_queue.queue_messages
acme_queue.outbox
acme_queue.outbox_consumer_groups
acme_queue.dead_letter_queue
acme_queue.queue_template

-- Event sourcing in another schema
CREATE SCHEMA acme_events;
acme_events.bitemporal_event_log
acme_events.event_store_template
```

**Configuration:**

```
flyway.schemas=acme_queue,acme_events
peegeq.db.schema=acme_queue
# Event store uses acme_events automatically
```

**Example 3: Multi-Tenant Architecture**

```
-- Tenant A
CREATE SCHEMA tenant_a_queue;
CREATE SCHEMA tenant_a_events;

-- Tenant B
CREATE SCHEMA tenant_b_queue;
CREATE SCHEMA tenant_b_events;

-- Tenant C
```

```
CREATE SCHEMA tenant_c_queue;
CREATE SCHEMA tenant_c_events;
```

## Configuration (per tenant):

```
// Tenant A configuration
PeeGeeQConfiguration tenantAConfig = new PeeGeeQConfiguration(
    "tenant_a",
    "localhost", 5432, "shared_db",
    "tenant_a_user", "password",
    "tenant_a_queue" // Tenant-specific schema
);

// Tenant B configuration
PeeGeeQConfiguration tenantBConfig = new PeeGeeQConfiguration(
    "tenant_b",
    "localhost", 5432, "shared_db",
    "tenant_b_user", "password",
    "tenant_b_queue" // Tenant-specific schema
);
```

## Example 4: Integration with Existing Database

```
-- Existing application schemas
CREATE SCHEMA app_core;      -- Your existing tables
CREATE SCHEMA app_reporting; -- Your existing reports

-- Add PeeGeeQ to existing database
CREATE SCHEMA app.messaging; -- PeeGeeQ queue operations
CREATE SCHEMA app.events;   -- PeeGeeQ event sourcing

-- All schemas coexist in same database
app_core.users
app_core.orders
app_reporting.sales_summary
app.messaging.queue_messages      -- PeeGeeQ
app.messaging.outbox              -- PeeGeeQ
app_events.bitemporal_event_log  -- PeeGeeQ
```

## Benefits:

- No separate database required
- Can use foreign keys between app and PeeGeeQ tables
- Shared connection pool
- Single backup/restore process

## Schema Migration Script Updates

When using custom schemas, Flyway automatically adjusts the migration scripts:

### Original Migration (uses `peegueq` schema):

```
-- V001__Create_Base_Tables.sql
CREATE SCHEMA IF NOT EXISTS peegueq;
CREATE TABLE peegueq.queue_messages (
    id BIGSERIAL PRIMARY KEY,
    --
    ...
);
```

### How Flyway Handles Custom Schemas:

Flyway's `schemas` configuration tells it which schema(s) to manage, and it automatically:

1. Creates the schema(s) if they don't exist
2. Sets the PostgreSQL `search_path` for the migration
3. Executes all `CREATE TABLE` statements in the configured schema(s)

**Note:** The migration script itself still contains `peegueq.` prefixes. If you need to use completely custom schema names in the SQL itself, you'll need to either:

#### 1. Use Flyway placeholders:

```
-- V001__Create_Base_Tables.sql
CREATE SCHEMA IF NOT EXISTS ${queueSchema};
CREATE TABLE ${queueSchema}.queue_messages (
    id BIGSERIAL PRIMARY KEY,
    --
    ...
);
```

```
# flyway.conf
flyway.placeholders.queueSchema=myapp.messaging
flyway.placeholders.eventSchema=myapp.events
```

2. **Create schema-specific migration files:** Duplicate V001 for different schema names (not recommended)

3. **Use dynamic SQL in migrations:** Use DO blocks to build table names dynamically (complex)

**Recommended Approach:** Use the standard `peegueq` and `bitemporal` schema names in migration files, and let Flyway's `schemas` configuration control where they're created. This keeps migrations simple and portable.

### Schema Permissions

When using custom schemas, ensure proper permissions:

```
-- Grant schema usage
GRANT USAGE ON SCHEMA myapp.messaging TO myapp_user;
GRANT USAGE ON SCHEMA myapp.events TO myapp_user;

-- Grant table permissions
GRANT SELECT, INSERT, UPDATE, DELETE ON ALL TABLES IN SCHEMA myapp.messaging TO
myapp_user;
GRANT SELECT, INSERT, UPDATE, DELETE ON ALL TABLES IN SCHEMA myapp.events TO
myapp_user;

-- Grant sequence permissions (for BIGSERIAL columns)
GRANT USAGE, SELECT ON ALL SEQUENCES IN SCHEMA myapp.messaging TO myapp_user;
GRANT USAGE, SELECT ON ALL SEQUENCES IN SCHEMA myapp.events TO myapp_user;

-- Grant future objects (for schema evolution)
ALTER DEFAULT PRIVILEGES IN SCHEMA myapp.messaging
    GRANT SELECT, INSERT, UPDATE, DELETE ON TABLES TO myapp_user;
ALTER DEFAULT PRIVILEGES IN SCHEMA myapp.messaging
    GRANT USAGE, SELECT ON SEQUENCES TO myapp_user;
```

## Search Path Considerations

### PostgreSQL Search Path:

By default, PostgreSQL uses `search_path = "$user", public`. When using custom schemas:

```
-- Option 1: Set search path for user
ALTER USER myapp_user SET search_path = myapp.messaging, myapp.events, public;

-- Option 2: Set search path for database
ALTER DATABASE myapp_db SET search_path = myapp.messaging, myapp.events, public;

-- Option 3: Set search path in session
SET search_path = myapp.messaging, myapp.events, public;
```

**PeeGeeQ Recommendation:** Always use **fully qualified table names** (schema.table) in application code to avoid `search_path` issues:

```
// ✓ GOOD: Fully qualified
String sql = "SELECT * FROM myapp.messaging.queue_messages WHERE topic = ?";

// ✗ BAD: Relies on search_path
String sql = "SELECT * FROM queue_messages WHERE topic = ?";
```

PeeGeeQ's internal code always uses fully qualified names, so `search_path` doesn't affect operations.

## Schema Best Practices

## 1. Use Descriptive Schema Names

GOOD:

```
myapp_queue, myapp_events  
tenant_acme.messaging, tenant_acme.events  
prod_peegeeq, prod_bitemporal
```

BAD:

```
schema1, schema2  
s1, s2  
temp_schema
```

## 2. Consistent Naming Convention

GOOD:

```
All tenants: tenant_{id}_queue, tenant_{id}_events  
All apps: {appname}_queue, {appname}_events
```

BAD:

```
Mixed: tenant_a_queue, tenantB.messaging, tenant-c-events
```

## 3. Document Schema Purpose

```
-- Add comments to schemas  
COMMENT ON SCHEMA myapp_queue IS 'PeeGeeQ queue operations for MyApp';  
COMMENT ON SCHEMA myapp_events IS 'PeeGeeQ event sourcing for MyApp';
```

## 4. Plan for Growth

GOOD: Leave room for expansion

```
myapp_queue (can add myapp_analytics, myapp_cache later)
```

BAD: Too generic

```
myapp (what happens when you need another schema?)
```

## 5. Schema Naming Length Limit

PostgreSQL identifiers (including schema names) have a 63-character limit.

GOOD:

```
tenant_acme.corp.queue (23 chars)
```

 BAD:

```
tenant_acme_corporation_very_long_name_queue.messaging_system (60+ chars)
```

## Schema Definition Sources and Alignment

Critical Understanding: Three Sources of Truth

PeeGeeQ maintains schema definitions in **three distinct locations**, each serving a specific purpose:

### PeeGeeQ Schema Definition Sources

#### 1. FLYWAY MIGRATIONS (Production Foundation)

Location: peegeeq-migrations/src/main/resources/db/migration/

Files:

```

    └── V001__Create_Base_Tables.sql (576 lines)
        └── Creates: outbox, queue_messages, dead_letter_queue,
                    bitemporal_event_log, message_processing, etc.

    └── V010__Create_Consumer_Group_Fanout_Tables.sql (442 lines)
        └── Adds: outbox_topics, outbox_topic_subscriptions, fanout columns

```

#### 2. SQL TEMPLATES (Dynamic Queue/EventStore Creation)

Location: peegeeq-db/src/main/resources/db/templates/

Structure:

```

    ├── base/          (36 files) - Schemas, extensions, template tables
    │   ├── 05-queue-template.sql      (queue_template table)
    │   └── 06-event-store-template.sql (event_store_template table)

    ├── queue/         (8 files) - Per-queue table creation
    │   ├── 01-table.sql           (LIKE queue_template INCLUDING ALL)
    │   ├── 02a-index-topic.sql
    │   ├── 02b-index-lock.sql
    │   └── ...

    └── eventstore/   (11 files) - Per-event-store table creation
        ├── 01-table.sql           (LIKE event_store_template INCLUDING ALL)
        ├── 02a-index-validtime.sql
        └── ...

```

#### 3. JAVA EMBEDDED SQL (Unit Test Fast Path)

Location: peegeeq-test-

support/src/main/java/.../PeeGeeQTestSchemaInitializer.java

Purpose: Unit tests that don't need full integration test isolation

Content: 747 lines of embedded CREATE TABLE statements

## Why Three Sources?

### 1. Flyway Migrations: Production Schema Foundation

**Purpose:** Version-controlled, auditable schema evolution for production databases

## What It Creates:

- Base tables that all queues share: `outbox`, `queue_messages`, `dead_letter_queue`
- Bitemporal event log infrastructure: `bitemporal_event_log`
- Consumer group fanout tables: `outbox_topics`, `outbox_topic_subscriptions`
- Template tables: `queue_template`, `event_store_template` (for PostgreSQL LIKE)
- Extensions: `uuid-ossp`, `pg_stat_statements`
- Schemas: `peegeq`, `bitemporal`

## What It Does NOT Create:

- Individual queue tables (e.g., `orders_queue`, `payments_queue`)
- Individual event store tables (e.g., `order_events`, `payment_events`)

## Used By:

- Production deployments (CI/CD pipelines)
- Development docker-compose setups
- Some integration tests that use Flyway directly

## File Example:

```
-- V001__Create_Base_Tables.sql (excerpt)
CREATE TABLE IF NOT EXISTS outbox (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    payload JSONB NOT NULL,
    -- ... all columns
);

CREATE TABLE IF NOT EXISTS queue_messages (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    -- ... all columns
);

-- Template table for dynamic queue creation
CREATE TABLE peegeq.queue_template (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    -- ... same columns as queue_messages
);
```

## 2. SQL Templates: Dynamic Table Creation Engine

**Purpose:** Parameterized SQL for creating tenant/feature-specific queue and event store tables at runtime

## What It Creates:

- Per-queue tables: `tenant_acme_orders`, `tenant_acme_payments`, `feature_x_queue`

- Per-event-store tables: `tenant_acme_events`, `order_event_log`, `payment_event_log`
- Dedicated indexes for each queue/event store
- Notification triggers for each queue/event store

**Key Feature:** Uses PostgreSQL `LIKE ... INCLUDING ALL` to inherit from template tables

### Used By:

- Production: `setupService.addQueue(setupId, config)` - creates new tenant queue
- Integration tests: `PeeGeeQDatabaseSetupService` - creates test-specific databases with multiple queues

### Template Example:

```
-- queue/01-table.sql
CREATE TABLE IF NOT EXISTS {schema}.{queueName} (
    LIKE peegeeq.queue_template INCLUDING ALL
);

-- queue/02a-index-topic.sql
CREATE INDEX idx_{queueName}_topic
ON {schema}.{queueName}(topic, visible_at, status);
```

### Processing Flow:

```
// Runtime: Parameters injected
Map<String, String> params = Map.of(
    "queueName", "tenant_acme_orders",
    "schema", "peegeeq"
);

// SqlTemplateProcessor replaces placeholders
// Result: CREATE TABLE peegeeq.tenant_acme_orders (
//           LIKE peegeeq.queue_template INCLUDING ALL
// );
```

## 3. Java Embedded SQL: Unit Test Fast Path

**Purpose:** Minimal schema setup for unit tests that need database but not full isolation

### What It Creates:

- Same base tables as Flyway migrations
- Embedded directly in Java code (no external file loading)
- Uses JDBC `Statement.execute()` directly

### Why Not Use Flyway or Templates?

- **Speed:** Faster than Flyway migration scanning

- **Simplicity:** No need for Vert.x or template processing
- **Isolation:** Unit tests don't need dynamic databases

## Used By:

- `peegeeq-outbox` unit tests
- `peegeeq-native` unit tests
- Tests that don't require test-per-class database isolation

## Code Example:

```
// PeeGeeQTestSchemaInitializer.java (excerpt)
private static void initializeOutboxSchema(Statement stmt) throws Exception {
    stmt.execute("""
        CREATE TABLE IF NOT EXISTS outbox (
            id BIGSERIAL PRIMARY KEY,
            topic VARCHAR(255) NOT NULL,
            payload JSONB NOT NULL,
            -- ... all columns (exact match to V001)
        )
    """);

    stmt.execute("CREATE INDEX IF NOT EXISTS idx_outbox_status_created " +
                "ON outbox(status, created_at)");
}
```

## Schema Alignment Strategy

### The Alignment Challenge

With three sources of schema definitions, **how do we prevent drift?**

### Example of Potential Drift:

```
-- Flyway V001: outbox table has 15 columns
ALTER TABLE outbox ADD COLUMN priority INT DEFAULT 5;

-- Template: queue_template not updated ✗
-- Unit Test SQL: PeeGeeQTestSchemaInitializer not updated ✗

Result: Production has priority column, tests don't → bugs!
```

## Alignment Mechanisms

### 1. Template Inheritance from Flyway (Automatic)

Templates use `LIKE ... INCLUDING ALL` to inherit from template tables created by Flyway:

```
-- Flyway V001 creates template table
CREATE TABLE peegeeq.queue_template (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    -- ... 15 columns
);

-- Template inherits ALL columns automatically
CREATE TABLE {schema}.{queueName} (
    LIKE peegeeq.queue_template INCLUDING ALL --  Automatic sync!
);
```

**Result:** When Flyway adds column to `queue_template`, ALL dynamically created queues get it automatically!

## 2. Contract Tests (Automated Validation)

**Location:** `peegeeq-`

`migrations/src/test/java/dev/mars/peegeeq/migrations/SchemaContractTest.java`

**Purpose:** Validate Flyway schema matches expectations

```
@Test
void testOutboxTableStructure() {
    // Connects to database created by Flyway
    ResultSet rs = connection.getMetaData().getColumns(null, null, "outbox",
null);

    // Validates ALL columns present
    Set<String> columns = extractColumnNames(rs);
    assertThat(columns).contains(
        "id", "topic", "payload", "status", "retry_count",
        "required_consumer_groups", // Added in V010
        "completed_consumer_groups", // Added in V010
        "completed_groups_bitmap" // Added in V010
    );
}

@Test
void testQueueMessagesTableStructure() {
    // Validates queue_messages table columns
    // Ensures template table matches
}
```

## 3. Manual Synchronization (Unit Test SQL)

**CRITICAL:** When Flyway migrations change base tables, `PeeGeeQTestSchemaInitializer.java` must be updated manually.

**Process:**

1. Developer adds column to Flyway migration (e.g., V011\_\_Add\_Priority.sql)
2. Developer updates `PeeGeeQTestSchemaInitializer.java` to match
3. Contract tests validate both match

### **Verification:**

```
# Run schema contract tests after migration changes
cd peegeeq-migrations
mvn test -Dtest=SchemaContractTest

# Run unit tests with updated schema
cd ../peegeeq-outbox
mvn test
```

### **4. Schema Drift Analysis (Manual Audit)**

**Location:** `peegeeq-migrations/docs/old/SCHEMA_DRIFT_ANALYSIS.md`

**Purpose:** Periodic manual review of schema alignment

**Results** (as of November 2025):

- **NO CRITICAL DRIFT** detected
- All Flyway migrations match Java code
- All template tables match base tables
- Unit test SQL matches Flyway migrations

### **Process:**

1. Quarterly schema audit
2. Compare Flyway SQL vs. Java embedded SQL
3. Compare template definitions vs. Flyway template tables
4. Document any intentional differences (future features)

## Alignment Best Practices

### **When Adding a New Column**

#### **Step-by-Step Process:**

1. Add to Flyway Migration  
File: `peegeeq-migrations/src/main/resources/db/migration/V011__Add_New_Column.sql`

```
ALTER TABLE outbox ADD COLUMN new_field VARCHAR(255);
ALTER TABLE queue_messages ADD COLUMN new_field VARCHAR(255);
```
2. Update Template Table (if needed)  
File: `peegeeq-`

```
migrations/src/main/resources/db/migration/V011__Add_New_Column.sql
```

```
ALTER TABLE peegeeq.queue_template ADD COLUMN new_field VARCHAR(255);

(Dynamic queues will inherit automatically via LIKE)
```

### 3. Update Unit Test SQL

File: peegeeq-test-support/src/main/java/.../PeeGeeQTestSchemaInitializer.java

```
stmt.execute("""
    CREATE TABLE IF NOT EXISTS outbox (
        id BIGSERIAL PRIMARY KEY,
        -- ... existing columns
        new_field VARCHAR(255) --  ADD HERE
    )
""");
```

### 4. Update Java Code

File: peegeeq-outbox/src/main/java/.../OutboxMessage.java

```
public class OutboxMessage {
    private String newField; //  ADD HERE
    // getters/setters
}
```

### 5. Run Contract Tests

```
cd peegeeq-migrations
mvn test -Dtest=SchemaContractTest
```

(Tests will fail if schema doesn't match expectations)

### 6. Update All Tests

```
cd ..
mvn clean test
```

(Ensure all tests pass with new schema)

## When Adding a New Table

**Critical:** Tables that ALL queues use (like `outbox`, `dead_letter_queue`) go in Flyway migrations. Tables created per-queue (like `tenant_X_orders`) are created by templates.

If table is shared across all queues:

- Add to Flyway migration (`V012__Add_Audit_Log.sql`)
- Add to `PeeGeeQTestSchemaInitializer.java`
- Add contract test to `SchemaContractTest.java`

If table is per-queue or per-tenant:

- Add template to `peegeeq-db/src/main/resources/db/templates/`
- Template inherits from existing template table
- No unit test update needed (uses template system)

## Schema Version Tracking

All three sources track schema versions differently:

1. Flyway Migrations
  - Table: flyway\_schema\_history
  - Tracks: V001, V010, V011, etc.
  - Location: Automatically created by Flyway
2. SQL Templates
  - Version: Embedded in .manifest files
  - Tracks: Execution order (01-table.sql, 02a-index-topic.sql, etc.)
  - Location: peegeeq-db/src/main/resources/db/templates/\*.manifest
3. Unit Test SQL
  - Table: schema\_version (manual)
  - Tracks: Version string (optional)
  - Location: Created by PeeGeeQTestSchemaInitializer

## Documentation Files

File	Purpose	Update Frequency
V001__Create_Base_Tables.sql	Production schema foundation	On breaking changes
V010__Create_Consumer_Group_Fanout_Tables.sql	Consumer group fanout feature	Feature additions
PEEGEEQ_COMPLETE_SCHEMA_SETUP.sql	Standalone setup script (documentation)	After each migration
PEEGEEQ_SCHEMA_QUICK_REFERENCE.sql	Quick reference (documentation)	After each migration
PeeGeeQTestSchemaInitializer.java	Unit test schema	Every migration change
SchemaContractTest.java	Automated validation	Every migration change
SCHEMA_DRIFT_ANALYSIS.md	Manual audit report	Quarterly

## Critical Takeaways

1. **Flyway Migrations = Foundation:** Creates schemas, extensions, base tables, and template tables
2. **SQL Templates = Dynamic Creation:** Creates tenant/feature-specific queues and event stores at runtime

3. **Unit Test SQL = Fast Path:** Minimal setup for unit tests without full integration test overhead
4. **Template Inheritance = Automatic Sync:** Dynamic queues inherit from template tables (no manual sync needed)
5. **Contract Tests = Validation:** Automated tests catch schema drift
6. **Manual Sync Required:** `PeeGeeQTestSchemaInitializer.java` must be updated when base tables change
7. **Quarterly Audits:** Manual schema drift analysis ensures long-term alignment

## Common Questions

### **Q: Why not use Flyway everywhere?**

A: Flyway is designed for schema evolution, not dynamic table creation. Templates enable runtime creation of tenant-specific queues.

### **Q: Why not use templates in unit tests?**

A: Templates require Vert.x and async processing. Unit tests use synchronous JDBC for simplicity and speed.

### **Q: What if I forget to update PeeGeeQTestSchemaInitializer?**

A: Unit tests will fail with "column does not exist" errors. Contract tests may also catch the issue.

### **Q: How do I know which source to update?**

A: Use the decision matrix above: Shared tables → Flyway + Unit Test SQL. Per-queue tables → Templates only.

### **Q: Are template tables used in production?**

A: Yes! Flyway creates `queue_template` and `event_store_template` tables. Runtime code uses `LIKE queue_template` to create tenant queues.

---

## Why Templates? Understanding the Design

### The Core Design: Dynamic Queue and Event Store Creation

**PeeGeeQ's fundamental architecture** is built around **runtime creation of queue and event store tables**.

This is not a testing convenience - it's the core production feature that enables:

### **Primary Design Goal: Parameterized Schema Objects**

PeeGeeQ creates **dedicated database tables for each queue and event store at runtime**:

- **Queue tables:** Each queue gets its own table with dedicated indexes and triggers
  - `peegueq.orders_queue` (created for "orders" queue)
  - `peegueq.payments_queue` (created for "payments" queue)
  - `peegueq.tenant_123_orders` (created for tenant 123)
- **Event store tables:** Each event store gets its own table with temporal indexes
  - `bitemporal.order_event_log` (created for order events)
  - `bitemporal.payment_event_log` (created for payment events)
  - `bitemporal.tenant_456_events` (created for tenant 456)

## Why This Design?

1. **Performance Isolation:** Each queue/event store has dedicated indexes and storage
2. **Multi-Tenancy:** Create tenant-specific queues without table bloat
3. **Feature Enablement:** Dynamically add queues when features are activated
4. **Scalability:** Partition data across multiple tables for horizontal scaling
5. **Maintenance:** Drop/archive old queues without affecting active ones

## Production Example:

```
// Application creates queues dynamically at runtime
QueueConfig config = new QueueConfig.Builder()
    .queueName("tenant_acme_orders") // Dynamic name!
    .maxRetries(3)
    .build();

// PeeGeeQ uses templates to create:
//   - peegeeq.tenant_acme_orders table
//   - 5 dedicated indexes
//   - notification trigger
setupService.addQueue(setupId, config).get();
```

## Secondary Benefit: Test Validation

Integration tests need to **validate this core production feature** by:

- Creating multiple queues with different names
- Verifying each queue has correct schema structure
- Testing parallel queue creation (simulating multi-tenant onboarding)
- Validating dynamic database creation (for isolated test environments)

**Test isolation is a side benefit**, not the primary reason for templates. Tests use the same template-based creation that production uses.

## The Solution: SQL Templates with Parameter Substitution

**Templates enable PeeGeeQ's core architecture by providing reusable, parameterized SQL patterns:**

### Static Migration (Production)

```
CREATE TABLE orders_queue (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255),
    ...
);
CREATE INDEX idx_orders_topic
    ON orders_queue(topic);
```

### Template (Integration Tests)

```
CREATE TABLE {queueName} (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255),
    ...
) INHERITS ({schema}.queue_template);
CREATE INDEX idx_{queueName}_topic
    ON {queueName}(topic);
```

## At runtime (production or test), templates are processed:

```
// Production: Creating queue dynamically
Map<String, String> params = Map.of(
    "queueName", "tenant_acme_orders",
    "schema", "peegueq"
);

// Template: CREATE TABLE {queueName} (...)
// Result:   CREATE TABLE tenant_acme_orders (...)
```

## Why Not Just Use Flyway Migrations?

Requirement	Flyway Migrations	SQL Templates
<b>Dynamic queue creation</b>	✗ No	✓ Yes (core feature)
<b>Parameterized table names</b>	✗ No	✓ Yes (core feature)
<b>Runtime queue creation</b>	✗ No	✓ Yes (core feature)
<b>Multi-tenant queue isolation</b>	✗ Limited	✓ Yes (dedicated tables)
<b>Feature flag queues</b>	✗ Requires migration	✓ Instant creation
<b>Version control</b>	✓ Yes	⚠ Template versions
<b>Schema evolution</b>	✓ Yes	⚠ Via base template updates

### Key Insight:

- **Flyway migrations** = Production schema **foundation** (extensions, schemas, base tables, template tables)
- **SQL templates** = Production **dynamic queue/event store creation** (per-queue tables, per-tenant tables, per-feature tables)
- **Both are used in production:** Migrations create the foundation once, templates create queues/event stores on-demand

### Production Architecture:

1. Deployment: Run Flyway migrations
  - |— Create schemas (peegueq, bitemporal)
  - |— Create extensions (uuid-ossp, pg\_stat\_statements)
  - |— Create base tables (queue\_messages, outbox, dead\_letter\_queue)
  - |— Create template tables (queue\_template, event\_store\_template)
2. Runtime: Use templates for dynamic creation
  - |— New tenant onboards → Create tenant\_X\_orders queue
  - |— Feature enabled → Create feature\_Y\_queue
  - |— Workflow defined → Create workflow\_Z\_step\_1\_queue

## Template Architecture Overview

### Template-Based Dynamic Queue Creation (Production & Test Environments)

#### 1. Application Requests Dynamic Queue Creation

```
// Production: New tenant onboarding
QueueConfig config = new QueueConfig.Builder()
    .queueName("tenant_acme_orders")
    .maxRetries(3)
    .build();

setupService.addQueue(setupId, config).get();

// Test: Validate dynamic creation capability
QueueConfig testConfig = new QueueConfig.Builder()
    .queueName("test_orders_queue")
    .build();
```

↓

#### 2. PeeGeeQDatabaseSetupService

- Loads SQL templates from resources
- Substitutes parameters in templates
- Executes SQL sequentially via Vert.x

↓

#### 3. SqlTemplateProcessor

```
Template: "CREATE TABLE {queueName} (...)"
Params:   {queueName: "tenant_acme_orders"}
Result:   "CREATE TABLE tenant_acme_orders (...)"
```

```
Executes → connection.query(processedSql).execute()
```

↓

#### 4. Database Tables Created Dynamically

Production Database: peegeeq\_prod  
Tables Created:

- peegeeq.tenant\_acme\_orders (dedicated)
- peegeeq.tenant\_acme\_payments (dedicated)
- bitemporal.tenant\_acme\_events (dedicated)

Indexes (per table):

- idx\_tenant\_acme\_orders\_topic
  - idx\_tenant\_acme\_orders\_visible
  - idx\_tenant\_acme\_orders\_priority
- Result: Tenant fully isolated at database level

## Real-World Production Example: Multi-Tenant SaaS

**Scenario:** Multi-tenant order processing system where each tenant needs isolated queues.

**Without Templates** (limited approach):

```
// ✗ Problem: All tenants share same queue_messages table
// - No performance isolation
// - No data isolation
// - Difficult to scale per-tenant
// - Cannot drop tenant data independently

@RestController
public class OrderController {

    @PostMapping("/orders")
    public void createOrder(@RequestHeader("X-Tenant-Id") String tenantId,
                           @RequestBody Order order) {
        // All tenants use same table
        queue.send(order); // queue_messages table

        // Problem: Millions of rows from all tenants in one table
        // Problem: Indexes cover all tenants (slower queries)
        // Problem: Cannot optimize per-tenant
    }
}
```

**With Templates** (PeeGeeQ's solution):

```
// ✓ Solution: Each tenant gets dedicated queue table
@RestController
public class OrderController {

    @Autowired
    private TenantQueueService tenantQueueService;

    @PostMapping("/orders")
    public CompletableFuture<OrderResponse> createOrder(
        @RequestHeader("X-Tenant-Id") String tenantId,
        @RequestBody Order order) {

        // Get or create tenant-specific queue (uses templates)
    }
}
```

```

        return tenantQueueService.getOrCreateTenantQueue(tenantId)
            .thenCompose(queue -> queue.send(order))
            .thenApply(msg -> new OrderResponse(msg.getId()));
    }
}

@Service
public class TenantQueueService {

    private final Map<String, Queue<Order>> tenantQueues = new ConcurrentHashMap<>();
    public CompletableFuture<Queue<Order>> getOrCreateTenantQueue(String tenantId) {
        if (tenantQueues.containsKey(tenantId)) {
            return CompletableFuture.completedFuture(tenantQueues.get(tenantId));
        }

        // Use templates to create tenant-specific queue
        QueueConfig config = new QueueConfig.Builder()
            .queueName("tenant_" + tenantId + "_orders") // Dynamic name!
            .maxRetries(3)
            .build();

        return setupService.addQueue(setupId, config)
            .thenApply(v -> {
                // Templates created:
                //   - peegeeq.tenant_acme_orders table
                //   - 5 dedicated indexes (only for this tenant)
                //   - notification trigger
                Queue<Order> queue = queueFactory.createQueue(
                    "tenant_" + tenantId + "_orders",
                    Order.class
                );
                tenantQueues.put(tenantId, queue);
                return queue;
            });
    }
}

```

## Benefits in Production:

- ✓ **Performance Isolation:** Each tenant's queue has dedicated indexes
- ✓ **Data Isolation:** Tenant data in separate tables (security, compliance)
- ✓ **Independent Scaling:** Scale storage per tenant
- ✓ **Tenant Lifecycle:** Drop tenant's table when they churn
- ✓ **Query Performance:** Indexes only cover one tenant's data
- ✓ **Maintenance:** Vacuum/analyze per tenant independently

## How Tests Validate This Production Feature

Integration tests verify that dynamic queue creation works correctly:

```

@Test
void testDynamicQueueCreation_ValidatesProductionCapability() {
    // This test validates the production feature of creating queues dynamically

    // Create first queue (simulates first tenant onboarding)
    QueueConfig config1 = new QueueConfig.Builder()
        .queueName("tenant_1_orders")
        .build();
    setupService.addQueue(setupId, config1).get();

    // Create second queue (simulates second tenant onboarding)
    QueueConfig config2 = new QueueConfig.Builder()
        .queueName("tenant_2_orders")
        .build();
    setupService.addQueue(setupId, config2).get();

    // Verify both queues exist independently
    verifyTableExists("peegeeq", "tenant_1_orders");
    verifyTableExists("peegeeq", "tenant_2_orders");

    // Verify each has dedicated indexes
    verifyIndexExists("idx_tenant_1_orders_topic");
    verifyIndexExists("idx_tenant_2_orders_topic");

    // This proves the production capability works!
}

```

### ### Template Inheritance and Table Templates

**\*\*Advanced Feature: PostgreSQL Table Inheritance\*\***

Templates use PostgreSQL's table inheritance to ensure consistency:

```

```sql
-- Step 1: Create template table (base/03a-table-queue-template.sql)
CREATE TABLE peegeeq.queue_template (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    payload JSONB NOT NULL,
    status VARCHAR(50) DEFAULT 'AVAILABLE',
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    -- ... all standard queue columns
);

-- Step 2: Create specific queue (queue/01-create-table.sql)
CREATE TABLE {schema}.{queueName} (
    -- Inherits ALL columns from queue_template
) INHERITS ({schema}.queue_template);

```

```
-- Result: orders_queue has exact same structure as queue_template
-- Benefit: Schema consistency enforced at database level
```

## Why This Matters:

1. **Schema Consistency**: All queues have identical structure
2. **Centralized Updates**: Change template once, all queues benefit
3. **Type Safety**: Database enforces column types
4. **Performance**: Inherited indexes and constraints

## How Templates Work

### Template Directory Structure

Templates are organized as **directories** containing numbered SQL files and a **.manifest** file:

```
src/main/resources/db/templates/
  base/
    .manifest          # Base schema (extensions, schemas, templates)
    01a-extension-uuid.sql      # Execution order manifest
    01b-extension-pgstat.sql    # One statement per file
    02a-schema-peeqeq.sql
    02b-schema-bitemporal.sql
    03a-table-queue-template.sql
    03b-table-eventstore-template.sql
    04a-table-queue-messages.sql
    04b-table-outbox.sql
    04c-table-dlq.sql
    05a-index-queue-topic.sql
    ... (31 files total)
    09e-consumer-index-topic.sql

  queue/
    .manifest          # Per-queue tables (parameterized)
    01-create-table.sql      # CREATE TABLE {queueName}
    02a-index-topic.sql    # CREATE INDEX idx_{queueName}_topic
    02b-index-visible.sql
    02c-index-status.sql
    03-function-notify.sql
    04-create-trigger.sql
    ... (8 files total)

  eventstore/
    .manifest          # Per-event-store tables (parameterized)
    01-create-table.sql      # CREATE TABLE {tableName}
    02a-index-event-id.sql    # CREATE INDEX idx_{tableName}_event_id
    02b-index-valid-time.sql
    02c-index-transaction-time.sql
```

```

└── ... (13 files total)
└── 04-create-trigger.sql

```

## Why Directories Instead of Single Files?

### Historical Context - The Vert.x Bug:

Originally, templates were single SQL files with multiple statements:

```

-- OLD APPROACH (BROKEN): peegeeq-template.sql
CREATE EXTENSION "uuid-ossp";          -- Statement 1
CREATE SCHEMA peegeeq;                  -- Statement 2
CREATE TABLE queue_template (...);     -- Statement 3
CREATE INDEX idx_queue_topic ON ...;   -- Statement 4
-- ... 28 more statements

-- Problem: Vert.x PostgreSQL client's .query().execute()
-- only executes the FIRST statement, silently ignores the rest!

```

**Tests passed for months** because:

1. PostgreSQL base image pre-included extensions
2. Idempotent SQL (`CREATE SCHEMA IF NOT EXISTS`) masked failures
3. Test execution order caused accidental success
4. Only verified 4 of 30+ database objects

See: [VERTX\\_MULTISTATEMENT\\_SQL\\_BUG\\_ANALYSIS.md](#) for complete analysis.

**Solution:** Split into one statement per file, execute sequentially:

```

-- NEW APPROACH (WORKS): Multiple files
-- base/01a-extension-uuid.sql
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";

-- base/01b-extension-pgstat.sql
CREATE EXTENSION IF NOT EXISTS "pg_stat_statements";

-- base/02a-schema-peegueq.sql
CREATE SCHEMA IF NOT EXISTS peegeeq;

-- base/02b-schema-bitemporal.sql
CREATE SCHEMA IF NOT EXISTS bitemporal;

-- ... one statement per file

```

## The Manifest File

**Purpose:** Defines execution order of SQL files in a directory.

**Format:** Plain text, one filename per line, comments start with #:

```
# base/.manifest - Execution order for base template

# Step 1: Extensions (required for UUID and monitoring)
01a-extension-uuid.sql
01b-extension-pgstat.sql

# Step 2: Schemas
02a-schema-peegeeq.sql
02b-schema-bitemporal.sql

# Step 3: Template tables (inheritance base)
03a-table-queue-template.sql
03b-table-eventstore-template.sql

# Step 4: Core operational tables
04a-table-queue-messages.sql
04b-table-outbox.sql
04c-table-dlq.sql

# Step 5: Indexes for performance
05a-index-queue-topic.sql
05b-index-queue-visible.sql
...
```

## Loading Process:

```
// SqlTemplateProcessor.loadTemplateFiles()

1. Look for .manifest file in template directory
2. If found:
   - Read filenames line by line
   - Load each SQL file in order
   - Return List<String> of SQL content
3. If not found:
   - Try to load directory name as single file (backward compatibility)
```

## Template Parameter Substitution

**Syntax:** Parameters use {parameterName} placeholder syntax.

### Example 1: Queue Template

```
-- queue/01-create-table.sql (BEFORE substitution)
CREATE TABLE {schema}.{queueName} (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
```

```

payload JSONB NOT NULL,
visible_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
status VARCHAR(50) DEFAULT 'AVAILABLE',
-- ... more columns
) INHERITS ({schema}.queue_template);

```

```

// Application code
Map<String, String> params = Map.of(
    "schema", "peegeeq",
    "queueName", "orders_queue"
);

templateProcessor.applyTemplateReactive(connection, "queue", params);

```

```

-- Result (AFTER substitution)
CREATE TABLE peegeeq.orders_queue (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    payload JSONB NOT NULL,
    visible_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    status VARCHAR(50) DEFAULT 'AVAILABLE',
    -- ... more columns
) INHERITS (peegeeq.queue_template);

```

## Example 2: Queue Index Template

```

-- queue/02a-index-topic.sql (BEFORE)
CREATE INDEX idx_{queueName}_topic
ON {schema}.{queueName}(topic, visible_at, status);

```

```

-- Result (AFTER substitution with params above)
CREATE INDEX idx_orders_queue_topic
ON peegeeq.orders_queue(topic, visible_at, status);

```

## Example 3: Event Store Template

```

-- eventstore/01-create-table.sql (BEFORE)
CREATE TABLE {schema}.{tableName} (
    id BIGSERIAL PRIMARY KEY,
    event_id VARCHAR(255) NOT NULL,
    event_type VARCHAR(255) NOT NULL,

```

```

valid_time TIMESTAMP WITH TIME ZONE NOT NULL,
transaction_time TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
payload JSONB NOT NULL,
-- ... more columns
) INHERITS ({schema}.event_store_template);

-- Notification trigger (uses parameter)
CREATE TRIGGER {tableName}_notify_trigger
AFTER INSERT ON {schema}.{tableName}
FOR EACH ROW
EXECUTE FUNCTION notify_event_insert('{notificationPrefix}');

```

```

// Application code
Map<String, String> params = Map.of(
    "schema", "bitemporal",
    "tableName", "order_event_log",
    "notificationPrefix", "order_event_"
);

templateProcessor.applyTemplateReactive(connection, "eventstore", params);

```

```

-- Result (AFTER substitution)
CREATE TABLE bitemporal.order_event_log (
    id BIGSERIAL PRIMARY KEY,
    event_id VARCHAR(255) NOT NULL,
    event_type VARCHAR(255) NOT NULL,
    valid_time TIMESTAMP WITH TIME ZONE NOT NULL,
    transaction_time TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    payload JSONB NOT NULL,
    -- ... more columns
) INHERITS (bitemporal.event_store_template);

CREATE TRIGGER order_event_log_notify_trigger
AFTER INSERT ON bitemporal.order_event_log
FOR EACH ROW
EXECUTE FUNCTION notify_event_insert('order_event_');

```

## SqlTemplateProcessor Implementation

### Core Logic:

```

public class SqlTemplateProcessor {

    /**
     * Apply SQL template directory with parameter substitution.
     *
     * CRITICAL: Vert.x PostgreSQL client's .query().execute() only executes

```

```
* the FIRST statement in multi-statement SQL. Templates are organized as
* directories with numbered files to work around this limitation.
*/
public Future<Void> applyTemplateReactive(
    SqlConnection connection,
    String templateDir,
    Map<String, String> parameters) {

    // Step 1: Load all SQL files from directory (using .manifest)
    List<String> sqlFiles = loadTemplateFiles(templateDir);

    // Step 2: Execute files sequentially using Future composition
    Future<Void> chain = Future.succeededFuture();

    for (String sqlContent : sqlFiles) {
        // Step 3: Substitute parameters in SQL
        String processedSql = processTemplate(sqlContent, parameters);

        // Step 4: Chain execution (wait for previous to complete)
        chain = chain.compose(v ->
            connection.query(processedSql).execute()
                .mapEmpty()
        );
    }

    return chain;
}

/**
 * Simple string replacement for parameters.
 */
private String processTemplate(String template, Map<String, String>
parameters) {
    String result = template;
    for (Map.Entry<String, String> entry : parameters.entrySet()) {
        // Replace {paramName} with actual value
        result = result.replace("{" + entry.getKey() + "}", entry.getValue());
    }
    return result;
}

/**
 * Load SQL files from directory using .manifest file.
 */
private List<String> loadTemplateFiles(String templateDir) throws IOException
{
    // 1. Try to load .manifest file
    String manifestPath = "/db/templates/" + templateDir + "/.manifest";
    InputStream manifestStream = getClass().getResourceAsStream(manifestPath);

    if (manifestStream != null) {
        // Read manifest and load files in order
        List<String> fileNames = readManifest(manifestStream);
        List<String> sqlContents = new ArrayList<>();
    }
}
```

```

        for (String fileName : fileNames) {
            String filePath = "/db/templates/" + templateDir + "/" + fileName;
            String content = loadFile(filePath);
            sqlContents.add(content);
        }

        return sqlContents;
    } else {
        // No manifest - try single file (backward compatibility)
        String filePath = "/db/templates/" + templateDir;
        String content = loadFile(filePath);
        return List.of(content);
    }
}
}
}

```

## Template Execution Flow

### Complete Example: Creating Orders Queue

```

// Step 1: Test defines what it needs
QueueConfig queueConfig = new QueueConfig.Builder()
    .queueName("orders_queue")
    .maxRetries(3)
    .build();

DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
    .setupId("test_12345")
    .databaseConfig(dbConfig)
    .queues(List.of(queueConfig))
    .build();

// Step 2: Setup service creates database and applies templates
setupService.createCompleteSetup(request).get();

```

### Behind the Scenes:

PeeGeeQDatabaseSetupService.createCompleteSetup()



1. Create Database  
CREATE DATABASE test\_12345



2. Apply Base Template (31 files)

```
SqlTemplateProcessor.applyTemplateReactive(
    connection, "base", Map.of()
)
```

Executes:

- 01a-extension-uuid.sql
- 01b-extension-pgstat.sql
- 02a-schema-peegeeq.sql
- ... (31 files sequentially)



### 3. Apply Queue Template (8 files, parameterized)

```
SqlTemplateProcessor.applyTemplateReactive(
    connection,
    "queue",
    Map.of(
        "schema", "peegeeq",
        "queueName", "orders_queue"
    )
)
```

Executes (after parameter substitution):

- CREATE TABLE peegeeq.orders\_queue (...)
- CREATE INDEX idx\_orders\_queue\_topic (...)
- CREATE INDEX idx\_orders\_queue\_visible (...)
- ... (8 files sequentially)



### 4. Create PeeGeeQManager

- Connects to test\_12345 database
- Creates QueueFactory for orders\_queue
- Returns DatabaseSetupResult

## Template Best Practices

### 1. One Statement Per File

```
-- ✓ GOOD: 01-create-table.sql
CREATE TABLE {schema}.{queueName} (...);
```

```
-- ✓ GOOD: 02-create-index.sql
CREATE INDEX idx_{queueName}_topic ON {schema}.{queueName}(topic);
```

```
-- ✗ BAD: create-queue.sql
CREATE TABLE {schema}.{queueName} (...);
```

```
CREATE INDEX idx_{queueName}_topic ON {schema}.{queueName}(topic);
-- Only first statement executes due to Vert.x limitation!
```

## 2. Use Idempotent SQL

```
-- ✓ GOOD: Safe to run multiple times
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";
CREATE SCHEMA IF NOT EXISTS peegeeq;
CREATE TABLE IF NOT EXISTS {schema}.queue_template (...);

-- ✗ BAD: Fails on second execution
CREATE EXTENSION "uuid-ossp";
CREATE SCHEMA peegeeq;
CREATE TABLE {schema}.queue_template (...);
```

## 3. Meaningful File Naming

```
✓ GOOD:
01a-extension-uuid.sql      # Clear: UUID extension
01b-extension-pgstat.sql    # Clear: pg_stat_statements
02a-schema-peegeeq.sql     # Clear: peegeeq schema
03a-table-queue-template.sql # Clear: queue template table

✗ BAD:
file1.sql
file2.sql
setup.sql
```

## 4. Sequential Numbering with Grouping

```
✓ GOOD: Groups with sub-ordering
01a-extension-uuid.sql
01b-extension-pgstat.sql      # Same group (extensions)
02a-schema-peegeeq.sql
02b-schema-bitemporal.sql    # Same group (schemas)
03a-table-queue-template.sql
03b-table-eventstore-template.sql # Same group (templates)

✗ BAD: No grouping visible
01-extension-uuid.sql
02-extension-pgstat.sql
03-schema-peegeeq.sql
04-schema-bitemporal.sql
```

## 5. Document Parameters in Comments

```
-- queue/01-create-table.sql
-- Parameters:
--   {schema} - Schema name (e.g., "peegeeq")
--   {queueName} - Queue table name (e.g., "orders_queue")

CREATE TABLE {schema}.{queueName} (
    ...
) INHERITS ({schema}.queue_template);
```

# Architecture Summary

## Three Database Setup Approaches

### 1. Flyway Migrations (Production & Development)

- **Module:** `peegeeq-migrations`
- **Method:** SQL migration scripts with version tracking
- **Files:** `V001__Create_Base_Tables.sql` (576 lines, complete schema)
- **Usage:** CI/CD pipelines, docker-compose, Maven commands
- **Advantages:**
  - Version control for schema changes
  - Rollback capability
  - Audit trail
  - Industry standard

### 2. Template-Based Setup (Integration Tests)

- **Class:** `PeeGeeQDatabaseSetupService` (825 lines)
- **Method:** SQL templates with parameter substitution
- **Files:** 52 SQL files in 3 directories (`base/`, `queue/`, `eventstore/`)
- **Usage:** Integration tests creating dynamic test databases
- **Advantages:**
  - Dynamic database creation
  - Parameterized table names
  - Isolated test environments
  - Per-test database cleanup

### 3. Direct SQL Injection (Unit Tests)

- **Class:** `PeeGeeQTestSchemaInitializer` (747 lines)
- **Method:** JDBC Statement execution of SQL scripts
- **Files:** Inline SQL strings for each schema component
- **Usage:** Unit tests requiring minimal schema setup
- **Advantages:**

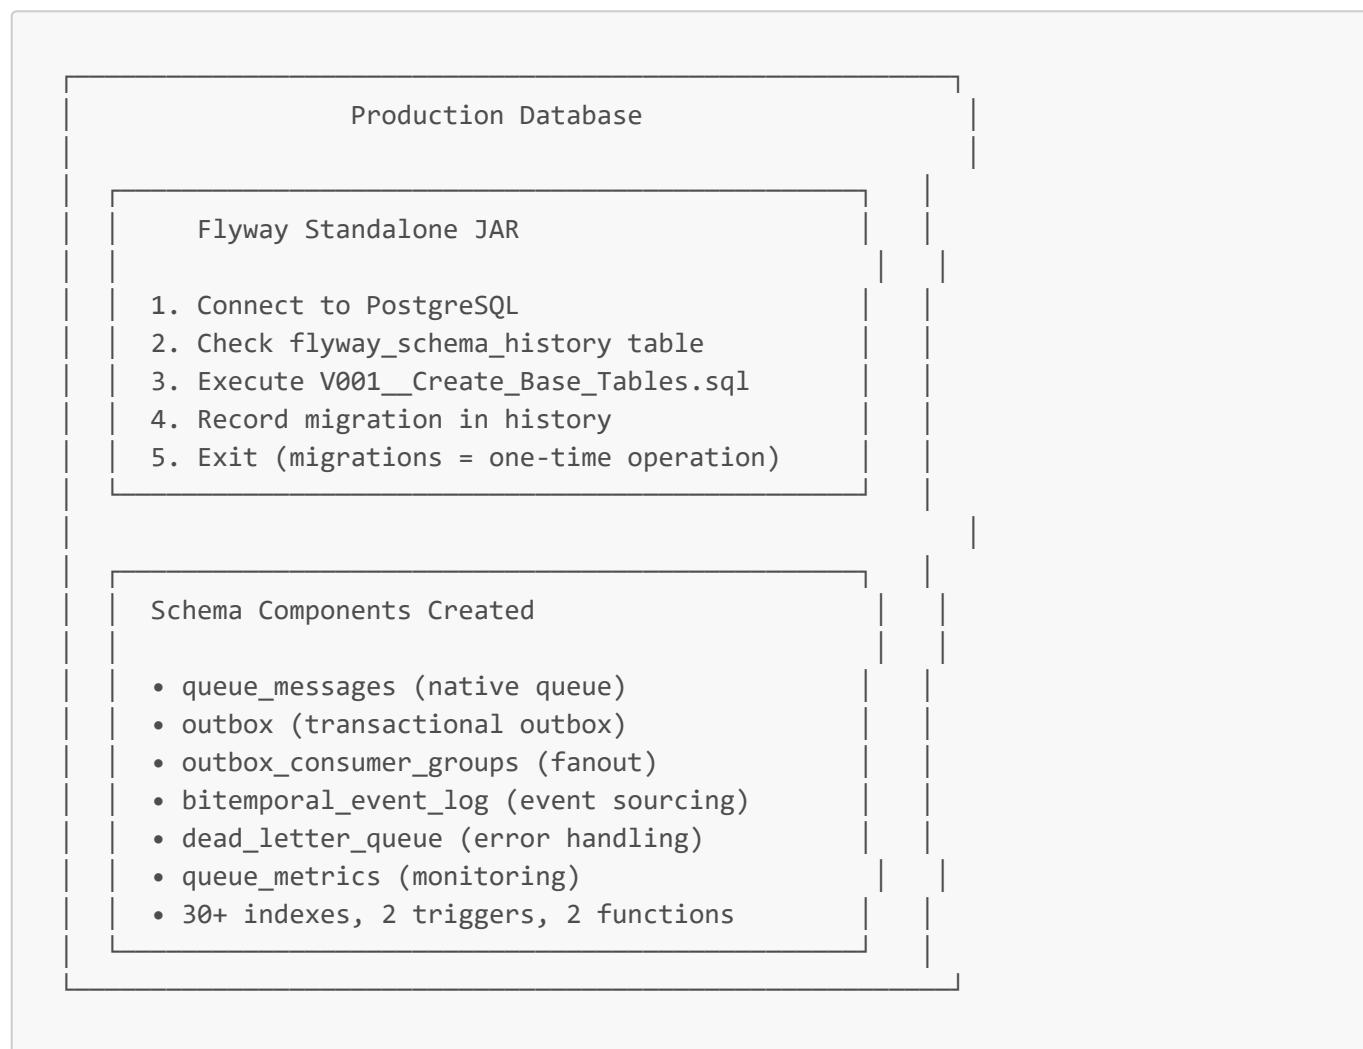
- Fast startup
  - Component-level granularity
  - No external dependencies
  - Simple and predictable
- 

## Production Setup

### Overview

Production databases are initialized using **Flyway migrations** from the `peegeq-migrations` module. This approach provides version control, audit trails, and rollback capabilities.

### Architecture



### Deployment Process

#### Step 1: Build Migration JAR

```
cd peegeq-migrations  
mvn clean package -DskipTests
```

Output: `target/peegeeq-migrations.jar` (standalone executable)

## Step 2: Run Migrations in CI/CD

### GitLab CI Example:

```
# .gitlab-ci.yml
stages:
  - build
  - migrate
  - deploy

migrate-production:
  stage: migrate
  script:
    - export DB_JDBC_URL=$PROD_DB_URL
    - export DB_USER=$PROD_DB_USER
    - export DB_PASSWORD=$PROD_DB_PASSWORD
    - java -jar peegeeq-migrations/target/peegeeq-migrations.jar migrate
  only:
    - main

deploy-application:
  stage: deploy
  dependencies:
    - migrate-production
  script:
    - # Deploy PeeGeeQ application
  only:
    - main
```

### GitHub Actions Example:

```
# .github/workflows/deploy.yml
name: Deploy Production

on:
  push:
    branches: [main]

jobs:
  migrate:
    name: Migrate Database
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - name: Setup Java
        uses: actions/setup-java@v3
        with:
          distribution: 'temurin'
          java-version: '23'
```

```

- name: Build Migration JAR
  run: |
    cd peegeeq-migrations
    mvn package -DskipTests
- name: Run Migrations
  env:
    DB_JDBC_URL: ${{ secrets.PROD_DB_URL }}
    DB_USER: ${{ secrets.PROD_DB_USER }}
    DB_PASSWORD: ${{ secrets.PROD_DB_PASSWORD }}
  run: |
    java -jar peegeeq-migrations/target/peegeeq-migrations.jar migrate

deploy:
  needs: migrate
  runs-on: ubuntu-latest
  steps:
    - name: Deploy Application
      run: |
        # Application deployment commands

```

### Step 3: Verify Migrations

```

# Check migration status
java -jar peegeeq-migrations.jar info \
  -url="jdbc:postgresql://prod-host:5432/peegeeq" \
  -user="produser" \
  -password="$PROD_PASSWORD"

# Expected output:
# +-----+-----+-----+-----+
# | Version | Description           | Installed on   | State   |
# +-----+-----+-----+-----+
# | 1       | Create Base Tables     | 2025-11-30 10:15:00 | Success |
# +-----+-----+-----+-----+

```

## Production Schema Components

When migrations complete successfully, the following components exist:

### Core Tables (7 tables)

- `schema_version` - Schema version tracking
- `queue_messages` - Native high-performance queue (10k+ msg/sec)
- `outbox` - Transactional outbox pattern (5k+ msg/sec)
- `outbox_consumer_groups` - Consumer group fanout
- `message_processing` - INSERT-only lock-free processing
- `dead_letter_queue` - Failed message handling
- `bitemporal_event_log` - Bi-temporal event sourcing (3k+ msg/sec)

## Monitoring Tables (2 tables)

- `queue_metrics` - Queue performance metrics
- `connection_pool_metrics` - Connection pool monitoring

## Indexes (30+ indexes)

- Performance indexes for all tables
- Composite indexes for bi-temporal queries
- GIN indexes for JSONB payload/header queries
- Unique indexes for constraint enforcement

## Functions & Triggers (4 objects)

- Trigger functions for automatic timestamp updates
- Notification triggers for real-time event streaming

## Views (3 views)

- `bitemporal_current_state` - Current event state
- `bitemporal_latest_events` - Most recent events
- `bitemporal_event_stats` - Event statistics

## Production Configuration

### Database Requirements

- **PostgreSQL:** 15.13 or higher
- **Extensions Required:**
  - `uuid-ossp` - UUID generation
  - `pg_stat_statements` - Query performance monitoring
- **Minimum Permissions:**
  - `CREATE EXTENSION` (for `uuid-ossp` and `pg_stat_statements`)
  - `CREATE SCHEMA` (for `peegeeq` and `bitemporal` schemas)
  - `CREATE TABLE, CREATE INDEX, CREATE TRIGGER`
  - `SELECT, INSERT, UPDATE, DELETE` on all tables

### Connection Configuration

```
# Production connection settings
flyway.url=jdbc:postgresql://prod-host:5432/peegeeq_prod
flyway.user=peegeeq_prod_user
flyway.password=${PROD_DB_PASSWORD}

# Schema Configuration (CUSTOMIZABLE)
# Default: peegeeq,bitemporal
# Custom examples:
#   - Single schema: myapp
```

```

#   - Tenant-specific: tenant_a_queue,tenant_a_events
#   - Namespace: acme.messaging,acme.events
flyway.schemas=peegeeq,bitemporal

flyway.locations=classpath:db/migration
flyway.baselineOnMigrate=true
flyway.validateOnMigrate=true

```

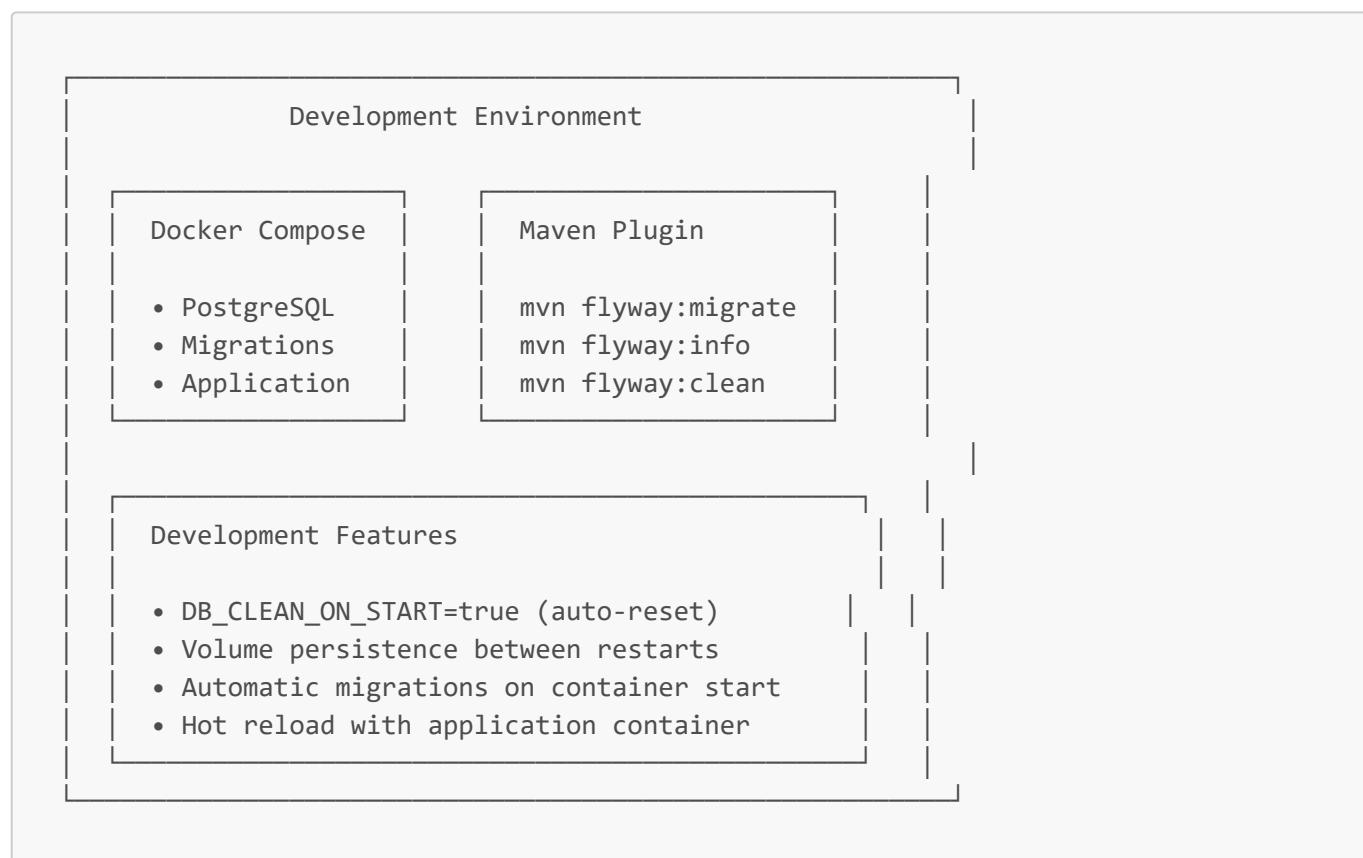
## Development Setup

### Overview

Development environments use **Flyway migrations** like production, but with additional tooling for rapid iteration:

- Docker Compose for automated local setup
- Maven plugin for manual migration control
- Development scripts for database reset

### Architecture



### Option 1: Docker Compose (Recommended)

**File:** `peegeeq-migrations/docker-compose.dev.yml`

### Architecture

```

services:
  postgres:
    image: postgres:15.13-alpine3.20
    ports: ["5432:5432"]
    environment:
      POSTGRES_DB: peegueq_dev
      POSTGRES_USER: peegueq_dev
      POSTGRES_PASSWORD: peegueq_dev

  migrations:
    build: ./Dockerfile
    depends_on:
      postgres: {condition: service_healthy}
    environment:
      DB_JDBC_URL: jdbc:postgresql://postgres:5432/peegueq_dev
      DB_USER: peegueq_dev
      DB_PASSWORD: peegueq_dev
      DB_CLEAN_ON_START: "true" # Clean database before migrations
    command: migrate
    restart: "no" # Run once and exit

```

## Usage

```

# Start PostgreSQL + run migrations
cd peegueq-migrations
docker-compose -f docker-compose.dev.yml up

# Migrations run automatically, then container exits
# PostgreSQL continues running with complete schema

# Reset database (clean + re-migrate)
docker-compose -f docker-compose.dev.yml down -v
docker-compose -f docker-compose.dev.yml up

```

## Advantages

- Automated setup (PostgreSQL + migrations)
- Clean environment every time
- Matches production container behavior
- No local Java/Maven required

## Option 2: Maven Plugin

**File:** `peegueq-migrations/pom.xml`

## Configuration

```

<plugin>
  <groupId>org.flywaydb</groupId>
  <artifactId>flyway-maven-plugin</artifactId>
  <version>10.23.0</version>
  <configuration>
    <url>jdbc:postgresql://localhost:5432/peegeeq_dev</url>
    <user>peegeeq_dev</user>
    <password>peegeeq_dev</password>
    <schemas>
      <schema>peegeeq</schema>
      <schema>bitemporal</schema>
    </schemas>
    <locations>
      <location>classpath:db/migration</location>
    </locations>
  </configuration>
</plugin>

```

## Usage

```

# Run migrations
cd peegeeq-migrations
mvn flyway:migrate

# Check migration status
mvn flyway:info

# Clean database (WARNING: deletes all data)
mvn flyway:clean

# Reset database (clean + migrate)
mvn flyway:clean flyway:migrate

# Validate migrations
mvn flyway:validate

```

## Advantages

- Fine-grained control
- Manual migration timing
- Easy to integrate with IDE
- No Docker required

## Option 3: Local Docker PostgreSQL

**File:** `docker-compose-local-dev.yml` (root directory)

## Configuration

```

services:
  postgres:
    image: postgres:15.13-alpine3.20
    ports: ["5432:5432"]
    environment:
      POSTGRES_DB: peegueq_dev
      POSTGRES_USER: peegueq_dev
      POSTGRES_PASSWORD: peegueq_dev
    volumes:
      - postgres_local_data:/var/lib/postgresql/data

```

## Usage

```

# Start PostgreSQL only
docker-compose -f docker-compose-local-dev.yml up -d

# Run migrations manually
cd peegueq-migrations
mvn flyway:migrate

# Application connects to localhost:5432

```

## Advantages

- Persistent data between sessions
- Controlled migration timing
- Faster restarts (no re-migration)
- IDE-friendly development

## Development Configuration

### Local Settings

```

# peegueq-migrations/src/main/resources/application-dev.properties
flyway.url=jdbc:postgresql://localhost:5432/peegueq_dev
flyway.user=peegueq_dev
flyway.password=peegueq_dev
flyway.cleanDisabled=false # Allow clean in dev (NEVER in prod)
flyway.baselineOnMigrate=true
flyway.outOfOrder=true

```

### Environment Variables (Alternative)

```
export DB_JDBC_URL="jdbc:postgresql://localhost:5432/peegeeq_dev"
export DB_USER="peegeeq_dev"
export DB_PASSWORD="peegeeq_dev"
export DB_CLEAN_ON_START="true" # Optional: clean before migrations
```

## Test Setup

### Overview

Test environments use **programmatic setup** instead of Flyway migrations. This allows:

- Dynamic database creation per test class
- Parameterized schema (e.g., unique table names)
- Fast cleanup and isolation
- No migration history pollution

### Integration Tests Architecture



## Integration Test Pattern

### Example Test Class

```
@TestInstance(TestInstance.Lifecycle.PER_CLASS)
public class QueueIntegrationTest {

    @Container
    private static PostgreSQLContainer<?> postgres = new PostgreSQLContainer<?>
("postgres:15.13-alpine3.20")
    .withDatabaseName("testdb")
    .withUsername("test")
    .withPassword("test");

    private PeeGeeQDatabaseSetupService setupService;
    private DatabaseSetupResult setupResult;

    @BeforeAll
    void setupDatabase() throws Exception {
        setupService = new PeeGeeQDatabaseSetupService();

        // Create dynamic test database
        DatabaseConfig dbConfig = new DatabaseConfig.Builder()
            .host(postgres.getHost())
            .port(postgres.getFirstMappedPort())
            .databaseName("test_db_" + UUID.randomUUID()) // Unique per test
            .username(postgres.getUsername())
            .password(postgres.getPassword())
            .schema("peegeeq")
            .build();

        // Define queues to create
        List<QueueConfig> queues = List.of(
            new QueueConfig.Builder()
                .queueName("test_queue_1")
                .maxRetries(3)
                .build()
        );

        // Define event stores to create
        List<EventStoreConfig> eventStores = List.of(
            new EventStoreConfig.Builder()
                .eventStoreName("test_events")
                .tableName("test_event_log")
                .notificationPrefix("test_event_")
                .build()
        );
    }
}
```

```

    // Create complete setup
    DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
        .setupId("test_setup_" + UUID.randomUUID())
        .databaseConfig(dbConfig)
        .queues(queues)
        .eventStores(eventStores)
        .build();

    setupResult = setupService.createCompleteSetup(request).get(30,
TimeUnit.SECONDS);

    // Verify setup successful
    assertEquals(DatabaseSetupStatus.READY, setupResult.getStatus());
}

@AfterAll
void cleanupDatabase() throws Exception {
    if (setupResult != null) {
        setupService.destroySetup(setupResult.getSetupId()).get(10,
TimeUnit.SECONDS);
    }
}

@Test
void testQueueOperations() {
    // Use setupResult.getQueueFactories() to get queue instances
    // Test queue send/receive operations
}
}

```

## Template Structure

Templates are organized as directories with manifest files:

```

src/main/resources/db/templates/
├── base/
│   ├── .manifest          # Lists files in order
│   ├── 01a-extension-uuid.sql      # CREATE EXTENSION "uuid-ossp"
│   ├── 01b-extension-pgstat.sql    # CREATE EXTENSION "pg_stat_statements"
│   ├── 02a-schema-peegueq.sql     # CREATE SCHEMA peegueq
│   ├── 02b-schema-bitemporal.sql  # CREATE SCHEMA bitemporal
│   ├── 03a-table-queue-template.sql # CREATE TABLE peegueq.queue_template
│   ├── ... (31 files total)
│   └── 09e-consumer-index-topic.sql

└── queue/
    ├── .manifest
    ├── 01-create-table.sql      # CREATE TABLE {schema}.{queueName}
    ├── 02a-index-topic.sql     # CREATE INDEX idx_{queueName}_topic
    ├── ... (8 files total)
    └── 04-create-trigger.sql

```

```

└── eventstore/
    ├── .manifest
    ├── 01-create-table.sql          # CREATE TABLE {schema}.{tableName}
    ├── 02a-index-event-id.sql     # CREATE INDEX idx_{tableName}_event_id
    ├── ... (13 files total)
    └── 04-create-trigger.sql

```

## Template Parameters

Templates support parameter substitution:

```

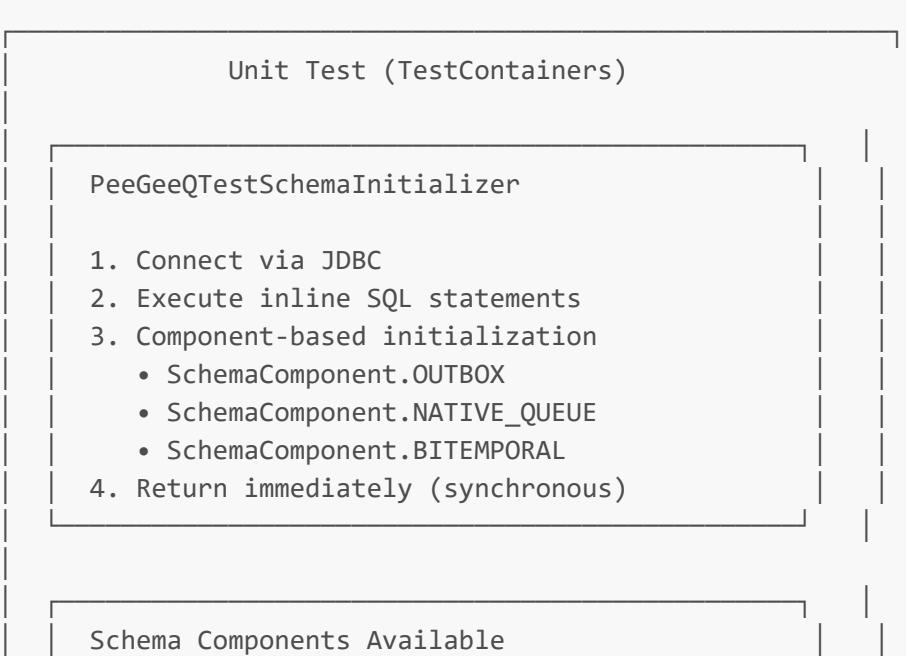
-- queue/01-create-table.sql
CREATE TABLE {schema}.{queueName} (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    payload JSONB NOT NULL,
    -- ... columns
) INHERITS ({schema}.queue_template);

-- Usage in code:
Map<String, String> params = Map.of(
    "schema", "peegeeq",
    "queueName", "orders_queue"
);
templateProcessor.applyTemplateReactive(connection, "queue", params);

-- Results in:
CREATE TABLE peegeeq.orders_queue (...) INHERITS (peegeeq.queue_template);

```

## Unit Tests Architecture



- SCHEMA\_VERSION
- OUTBOX
- NATIVE\_QUEUE
- DEAD\_LETTER\_QUEUE
- BITEMPORAL
- METRICS
- CONSUMER\_GROUP\_FANOUT
- QUEUE\_ALL (outbox + native + dlq)
- ALL (everything)

## Example Unit Test

```
@TestInstance(TestInstance.Lifecycle.PER_CLASS)
public class OutboxServiceTest {

    @Container
    private static PostgreSQLContainer<?> postgres = new PostgreSQLContainer<>(
        "postgres:15.13-alpine3.20");

    @BeforeAll
    void setupSchema() {
        // Initialize only what we need for this test
        PeeGeeQTestSchemaInitializer.initializeSchema(
            postgres,
            SchemaComponent.OUTBOX,
            SchemaComponent.DEAD_LETTER_QUEUE
        );
    }

    @AfterEach
    void cleanupTestData() {
        // Clean test data but keep schema
        PeeGeeQTestSchemaInitializer.cleanupTestData(
            postgres,
            SchemaComponent.OUTBOX
        );
    }

    @Test
    void testOutboxInsertion() {
        // Test outbox operations
    }
}
```

## Available Schema Components

```

public enum SchemaComponent {
    SCHEMA_VERSION, // Schema version tracking table
    OUTBOX, // Outbox pattern tables + indexes
    NATIVE_QUEUE, // Native queue tables + indexes
    DEAD_LETTER_QUEUE, // Dead letter queue table
    BITEMPORAL, // Bi-temporal event log + views + triggers
    METRICS, // Metrics and monitoring tables
    CONSUMER_GROUP_FANOUT, // Consumer group fanout tables
    QUEUE_ALL, // All queue-related (OUTBOX + NATIVE + DLQ)
    ALL // Everything
}

```

## Test Configuration Best Practices

### 1. Database Isolation

```

// Create unique database per test class
String dbName = "test_" + getClass().getSimpleName() + "_" + UUID.randomUUID();

```

### 2. Resource Cleanup

```

@AfterAll
void cleanup() {
    // Always cleanup in @AfterAll
    if (setupResult != null) {
        setupService.destroySetup(setupResult.getSetupId()).get();
    }
}

```

### 3. Parallel Test Execution

```

// Use unique setup IDs for parallel tests
String setupId = "test_" + Thread.currentThread().getId() + "_" +
UUID.randomUUID();

```

## Dynamic Queue Creation at Runtime

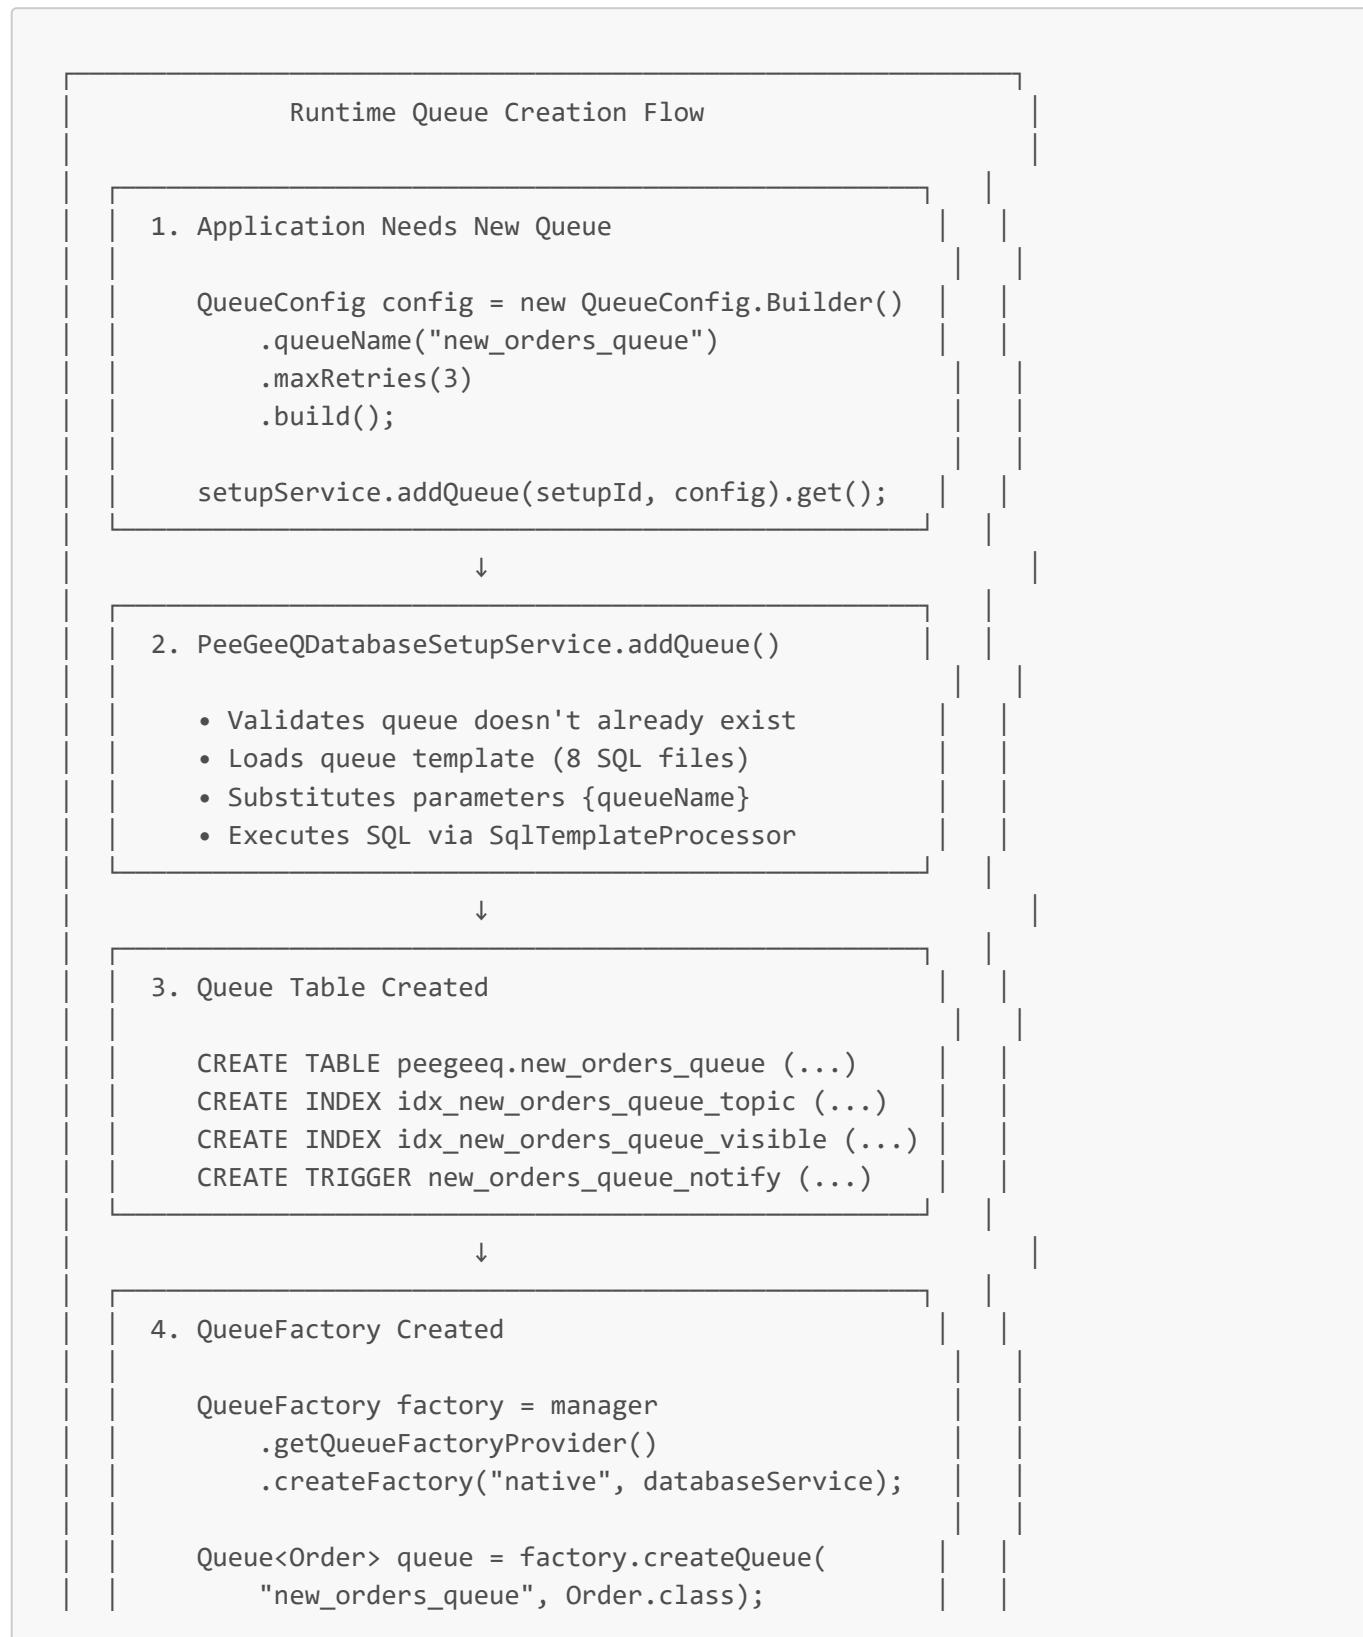
### Overview

PeeGeeQ supports **dynamic queue creation at runtime** - the ability to create new queue tables on-demand without restarting the application or running migrations. This is critical for:

- **Multi-tenant systems:** Create queues dynamically when onboarding new tenants
- **Feature flags:** Enable new queues when features are activated
- **Dynamic workflows:** Create queues based on user-defined business processes
- **Auto-scaling:** Add queues dynamically as load increases
- **Plugin architectures:** Third-party plugins can create their own queues

## How Dynamic Queue Creation Works

### Architecture



## 5. Queue Ready for Use

```
queue.send(new Order(...));
queue.receive().thenAccept(msg -> ...);
```

## Production Usage

### Example 1: Multi-Tenant Queue Creation

```
/*
 * Service that creates tenant-specific queues on demand
 */
public class TenantQueueService {

    private final PeeGeeQDatabaseSetupService setupService;
    private final String setupId;
    private final Map<String, QueueFactory> tenantQueues = new ConcurrentHashMap<>();
()

    public CompletableFuture<Queue<Order>> getOrCreateTenantQueue(String tenantId) {
        String queueName = "tenant_" + tenantId + "_orders";

        // Check if queue already exists
        if (tenantQueues.containsKey(queueName)) {
            return CompletableFuture.completedFuture(
                tenantQueues.get(queueName).createQueue(queueName, Order.class)
            );
        }

        // Create new queue dynamically
        QueueConfig config = new QueueConfig.Builder()
            .queueName(queueName)
            .maxRetries(3)
            .retryDelayMillis(5000)
            .build();

        return setupService.addQueue(setupId, config)
            .thenApply(v -> {
                // Create queue factory
                QueueFactory factory = createQueueFactory();
                tenantQueues.put(queueName, factory);

                logger.info("Created dynamic queue for tenant: {}", tenantId);
                return factory.createQueue(queueName, Order.class);
            });
    }
}
```

```

    }

}

// Usage in application
@RestController
@RequestMapping("/api/tenants")
public class TenantController {

    @Autowired
    private TenantQueueService tenantQueueService;

    @PostMapping("/{tenantId}/orders")
    public CompletableFuture<OrderResponse> createOrder(
        @PathVariable String tenantId,
        @RequestBody OrderRequest request) {

        // Get or create tenant-specific queue
        return tenantQueueService.getOrCreateTenantQueue(tenantId)
            .thenCompose(queue -> {
                Order order = new Order(tenantId, request);
                return queue.send(order);
            })
            .thenApply(msg -> new OrderResponse(msg.getId()));
    }
}
}

```

## Example 2: Feature Flag Activated Queue

```

/**
 * Creates queues dynamically when features are enabled
 */
public class FeatureQueueManager {

    private final PeeGeeQDatabaseSetupService setupService;
    private final String setupId;
    private final Set<String> activeQueues = ConcurrentHashMap.newKeySet();

    @EventListener
    public void onFeatureEnabled(FeatureEnabledEvent event) {
        String featureName = event.getFeatureName();
        String queueName = featureName.toLowerCase() + "_queue";

        if (activeQueues.contains(queueName)) {
            logger.info("Queue {} already exists", queueName);
            return;
        }

        logger.info("Creating queue for newly enabled feature: {}", featureName);

        QueueConfig config = new QueueConfig.Builder()
            .queueName(queueName)

```

```

        .maxRetries(5)
        .build();

    setupService.addQueue(setupId, config)
        .thenRun(() -> {
            activeQueues.add(queueName);
            logger.info("Queue {} created and ready", queueName);
        })
        .exceptionally(ex -> {
            logger.error("Failed to create queue for feature: {}", featureName, ex);
            return null;
        });
    }
}

```

### Example 3: Dynamic Workflow Queues

```

/**
 * Creates queues for user-defined workflows
 */
public class WorkflowQueueService {

    public CompletableFuture<Void> createWorkflowQueues(Workflow workflow) {
        List<CompletableFuture<Void>> queueCreations = new ArrayList<>();

        // Create a queue for each step in the workflow
        for (WorkflowStep step : workflow.getSteps()) {
            String queueName = String.format("workflow_%s_step_%s",
                workflow.getId(), step.getName());

            QueueConfig config = new QueueConfig.Builder()
                .queueName(queueName)
                .maxRetries(step.getMaxRetries())
                .build();

            CompletableFuture<Void> future = setupService.addQueue(setupId,
config);
            queueCreations.add(future);
        }

        // Wait for all queues to be created
        return CompletableFuture.allOf(
            queueCreations.toArray(new CompletableFuture[0]))
                .thenRun(() -> {
                    logger.info("Created {} queues for workflow: {}",
                        queueCreations.size(), workflow.getId());
                });
    }
}

```

## Development Usage

Dynamic queue creation is also used during development for:

### Hot Reload Support

```
/*
 * Development mode: Create queues on-the-fly without restart
 */
@Profile("development")
@Component
public class DevQueueManager {

    @PostMapping("/dev/queues/create")
    public CompletableFuture<String> createDevQueue(@RequestParam String name) {
        QueueConfig config = new QueueConfig.Builder()
            .queueName("dev_" + name)
            .maxRetries(1)
            .build();

        return setupService.addQueue(setupId, config)
            .thenApply(v -> "Queue created: dev_" + name);
    }
}
```

## Integration Test Usage

Dynamic queue creation is extensively used in integration tests:

```
@Test
void testDynamicQueueCreation() throws Exception {
    // Start with no queues
    DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
        .setupId("test_dynamic")
        .databaseConfig(dbConfig)
        .queues(List.of()) // No initial queues
        .build();

    DatabaseSetupResult result = setupService.createCompleteSetup(request).get();

    // Later, add queue dynamically
    QueueConfig queueConfig = new QueueConfig.Builder()
        .queueName("dynamic_test_queue")
        .maxRetries(3)
        .build();

    setupService.addQueue("test_dynamic", queueConfig).get();

    // Verify queue exists and works
    DatabaseConfig dbConfig = result.getDatabaseConfig();
```

```

try (Connection conn = DriverManager.getConnection(
    dbConfig.getJdbcUrl(),
    dbConfig.getUsername(),
    dbConfig.getPassword())) {

    ResultSet rs = conn.createStatement().executeQuery(
        "SELECT COUNT(*) FROM information_schema.tables " +
        "WHERE table_schema = 'peegeeq' " +
        "AND table_name = 'dynamic_test_queue'" +
    );

    rs.next();
    assertEquals(1, rs.getInt(1), "Queue table should exist");
}
}

```

## What Gets Created

When a queue is created dynamically, the following database objects are created:

### Queue Table

```

CREATE TABLE {schema}.{queueName} (
    id BIGSERIAL PRIMARY KEY,
    topic VARCHAR(255) NOT NULL,
    payload JSONB NOT NULL,
    visible_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    lock_id BIGINT,
    lock_until TIMESTAMP WITH TIME ZONE,
    retry_count INT DEFAULT 0,
    max_retries INT DEFAULT 3,
    status VARCHAR(50) DEFAULT 'AVAILABLE',
    headers JSONB DEFAULT '{}',
    error_message TEXT,
    correlation_id VARCHAR(255),
    message_group VARCHAR(255),
    priority INT DEFAULT 5
) INHERITS ({schema}.queue_template);

```

### Performance Indexes (5 indexes)

```

CREATE INDEX idx_{queueName}_topic
    ON {schema}.{queueName}(topic, visible_at, status);

CREATE INDEX idx_{queueName}_visible
    ON {schema}.{queueName}(visible_at) WHERE status = 'AVAILABLE';

```

```

CREATE INDEX idx_{queueName}_lock
    ON {schema}.{queueName}(lock_id) WHERE lock_id IS NOT NULL;

CREATE INDEX idx_{queueName}_status
    ON {schema}.{queueName}(status, created_at);

CREATE INDEX idx_{queueName}_priority
    ON {schema}.{queueName}(priority, created_at);

```

## Notification Trigger

```

CREATE TRIGGER {queueName}_notify_trigger
AFTER INSERT ON {schema}.{queueName}
FOR EACH ROW
EXECUTE FUNCTION notify_queue_insert();

```

## Function (shared across all queues)

```

CREATE OR REPLACE FUNCTION notify_queue_insert()
RETURNS TRIGGER AS $$ 
BEGIN
    PERFORM pg_notify('queue_insert', NEW.topic || ':' || NEW.id);
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

```

## Dynamic Event Store Creation

Similarly, event stores can be created dynamically:

```

/**
 * Create event store dynamically at runtime
 */
public CompletableFuture<EventStore<OrderEvent>> getOrCreateEventStore(String
storeName) {
    EventStoreConfig config = new EventStoreConfig.Builder()
        .eventStoreName(storeName)
        .tableName(storeName + "_log")
        .notificationPrefix(storeName + "_event_")
        .build();

    return setupService.addEventStore(setupId, config)
        .thenApply(v -> {
            // Create event store factory
            EventStoreFactory factory = new PeeGeeQEventStoreFactory(manager);
            return factory.createEventStore(

```

```

        OrderEvent.class,
        "bitemporal." + storeName + "_log"
    );
}
}
}

```

### Event Store Objects Created:

- Event log table with bi-temporal columns
- 10+ indexes for temporal queries
- GIN indexes for JSONB payload/headers
- Notification trigger for event streaming
- Inherited from `event_store_template`

## Performance Considerations

### 1. Creation Time

Dynamic queue creation takes approximately:

- Queue table: ~50ms
- 5 indexes: ~100ms
- Trigger: ~20ms
- Total: ~170ms per queue

### 2. Connection Pooling

```

// Use connection pool for dynamic creation
@Configuration
public class DynamicQueueConfig {

    @Bean
    public PeeGeeQDatabaseSetupService setupService() {
        // Setup service uses connection pooling internally
        return new PeeGeeQDatabaseSetupService();
    }
}

```

### 3. Concurrent Creation

```

// Safe for concurrent queue creation
CompletableFuture<Void> future1 = setupService.addQueue(setupId, config1);
CompletableFuture<Void> future2 = setupService.addQueue(setupId, config2);
CompletableFuture<Void> future3 = setupService.addQueue(setupId, config3);

```

```
// All three queues created in parallel (if different names)
CompletableFuture.allOf(future1, future2, future3).get();
```

## 4. Idempotent Creation

```
// Safe to call multiple times with same queue name
// Uses CREATE TABLE IF NOT EXISTS internally
setupService.addQueue(setupId, config).get(); // Creates table
setupService.addQueue(setupId, config).get(); // No-op (already exists)
```

## Limitations and Considerations

### 1. Schema Must Exist

```
// ✗ BAD: Schema doesn't exist
QueueConfig config = new QueueConfig.Builder()
    .queueName("orders_queue")
    .build();

// Schema "peegeeq" must be created first (via base template or migrations)
```

### 2. Template Table Required

```
// ✗ BAD: queue_template doesn't exist
// Dynamic queue creation requires queue_template to exist

// ✓ GOOD: Apply base template first
setupService.applySchemaTemplatesAsync(request).get(); // Creates templates
setupService.addQueue(setupId, queueConfig).get(); // Now works
```

### 3. Naming Constraints

```
// ✗ BAD: Invalid table names
"orders-queue"      // Hyphens not allowed
"123_orders"       // Can't start with number
"order queue"       // Spaces not allowed

// ✓ GOOD: Valid table names
"orders_queue"
"orders_v2"
"tenant_123_orders"
```

## 4. PostgreSQL Identifier Limit

Maximum table name length: 63 characters

✗ BAD:

```
"tenant_acme_corporation_very_long_company_name_orders_queue_v2" // 64 chars
```

✓ GOOD:

```
"tenant_acme_corp_orders_queue" // 31 chars
```

## Best Practices

### 1. Queue Naming Convention

```
// ✓ GOOD: Consistent naming
"tenant_{tenantId}_orders"           // Multi-tenant
"workflow_{workflowId}_step_{name}"   // Workflow steps
"feature_{featureName}_events"       // Feature-based

// ✗ BAD: Inconsistent naming
"orders_queue_for_tenant_123"
"tenant_456_queue_orders"
"payments_queue_tenant789"
```

### 2. Cache Queue Factories

```
// ✓ GOOD: Cache factories to avoid recreation
private final Map<String, QueueFactory> queueFactories = new ConcurrentHashMap<>();
();

public QueueFactory getQueueFactory(String queueName) {
    return queueFactories.computeIfAbsent(queueName, name ->
        manager.getQueueFactoryProvider()
            .createFactory("native", manager.getDatabaseService())
    );
}

// ✗ BAD: Create factory every time
public QueueFactory getQueueFactory(String queueName) {
    return manager.getQueueFactoryProvider()
        .createFactory("native", manager.getDatabaseService()); // Expensive!
}
```

### 3. Track Created Queues

```
// ✓ GOOD: Maintain registry of dynamically created queues
@Component
public class QueueRegistry {

    private final Set<String> createdQueues = ConcurrentHashMap.newKeySet();

    public CompletableFuture<Void> createQueueIfNotExists(String queueName) {
        if (createdQueues.contains(queueName)) {
            return CompletableFuture.completedFuture(null);
        }

        QueueConfig config = new QueueConfig.Builder()
            .queueName(queueName)
            .build();

        return setupService.addQueue(setupId, config)
            .thenRun(() -> {
                createdQueues.add(queueName);
                logger.info("Registered new queue: {}", queueName);
            });
    }

    public boolean queueExists(String queueName) {
        return createdQueues.contains(queueName);
    }

    public Set<String> getAllQueues() {
        return Collections.unmodifiableSet(createdQueues);
    }
}
```

#### 4. Handle Creation Failures Gracefully

```
// ✓ GOOD: Proper error handling
public CompletableFuture<Queue<Order>> getOrCreateQueue(String queueName) {
    return createQueueIfNeeded(queueName)
        .thenApply(v -> queueFactory.createQueue(queueName, Order.class))
        .exceptionally(ex -> {
            logger.error("Failed to create queue: {}", queueName, ex);

            // Check if queue already exists (race condition)
            if (queueExists(queueName)) {
                return queueFactory.createQueue(queueName, Order.class);
            }

            throw new QueueCreationException("Unable to create queue: " +
queueName, ex);
        });
}
```

## 5. Monitor Queue Creation

```
//  GOOD: Emit metrics for queue creation
@Component
public class MonitoredQueueService {

    @Autowired
    private MeterRegistry meterRegistry;

    public CompletableFuture<Void> createQueue(QueueConfig config) {
        Timer.Sample sample = Timer.start(meterRegistry);

        return setupService.addQueue(setupId, config)
            .whenComplete((result, error) -> {
                sample.stop(Timer.builder("peegeq.queue.creation")
                    .tag("queue", config.getQueueName())
                    .tag("success", error == null ? "true" : "false")
                    .register(meterRegistry));

                if (error == null) {
                    meterRegistry.counter("peegeq.queue.created").increment();
                } else {
                    meterRegistry.counter("peegeq.queue.creation.failed").increment();
                }
            });
    }
}
```

### Comparison: Static vs Dynamic Queue Creation

Aspect	Static (Migrations)	Dynamic (Runtime)
<b>When</b>	Application startup	On-demand
<b>Speed</b>	All queues created upfront	Only when needed
<b>Flexibility</b>	Fixed set of queues	Create as needed
<b>Migration</b>	Version controlled	No migration needed
<b>Restart Required</b>	Yes (for new queues)	No
<b>Multi-tenant</b>	All tenants upfront	Per-tenant on-demand
<b>Use Case</b>	Known queue set	Dynamic queue needs
<b>Production</b>	<input checked="" type="checkbox"/> Recommended for fixed queues	<input checked="" type="checkbox"/> Recommended for dynamic needs

### When to Use Each Approach

**Use Static Migration** when:

- Queue set is known at design time
- All applications need same queues
- Version control of schema is important
- Production stability is critical
- Schema changes need approval

**Use Dynamic Creation** when:

- Multi-tenant with tenant-specific queues
- Feature flags enable new queues
- User-defined workflows need queues
- Plugin architecture with plugin-specific queues
- A/B testing with experimental queues

**Use Both** (hybrid approach):

```
// Static: Core queues in migrations
// V001__Create_Base_Tables.sql creates:
//   - orders_queue
//   - payments_queue
//   - notifications_queue

// Dynamic: Tenant-specific queues at runtime
public CompletableFuture<Queue<Order>> getTenantQueue(String tenantId) {
    String queueName = "tenant_" + tenantId + "_orders";
    return getOrCreateQueue(queueName);
}
```

## Service-Specific Patterns

peegeeq-db (Core Database Module)

**Purpose:** Core database access and management

**Setup Method:** Uses all three approaches depending on context

- **Production:** Flyway migrations via `peegeeq-migrations`
- **Integration Tests:** `PeeGeeQDatabaseSetupService`
- **Unit Tests:** `PeeGeeQTestSchemaInitializer`

**Key Classes:**

- `PeeGeeQDatabaseSetupService` (825 lines)
  - Creates test databases dynamically
  - Applies SQL templates from `src/main/resources/db/templates/`
  - Manages PeeGeeQManager lifecycle
  - Handles cleanup and resource disposal

- **SqlTemplateProcessor** (153 lines)
  - Loads SQL files from template directories
  - Processes **.manifest** files for execution order
  - Substitutes parameters in SQL ({queueName}, {tableName})
  - Executes SQL using Vert.x reactive PostgreSQL client

### **Database Objects Created:**

- Core operational tables (queue\_messages, outbox, dead\_letter\_queue)
- Template tables (queue\_template, event\_store\_template)
- Consumer group tables (5 tables for fanout pattern)
- Indexes (30+ for performance)
- Functions and triggers (2 each)

### **Usage Example:**

```
// Integration test
PeeGeeQDatabaseSetupService setupService = new PeeGeeQDatabaseSetupService();
DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
    .setupId("test_123")
    .databaseConfig(dbConfig)
    .queues(List.of(queueConfig))
    .eventStores(List.of(eventStoreConfig))
    .build();
DatabaseSetupResult result = setupService.createCompleteSetup(request).get();
```

## peegeeq-native (Native Queue Implementation)

**Purpose:** High-performance native PostgreSQL queue (10k+ msg/sec)

### **Database Requirements:**

- **queue\_messages** table (from Flyway migrations or setup service)
- **message\_processing** table (INSERT-only lock-free processing)
- **dead\_letter\_queue** table (failed message handling)
- Indexes on topic, visible\_at, status, priority

### **Setup Patterns:**

#### **Production**

```
// Database already initialized via Flyway
// Native queue connects to existing schema
PeeGeeQConfiguration config = new PeeGeeQConfiguration();
PeeGeeQManager manager = new PeeGeeQManager(config);

// Get native queue factory
QueueFactory factory = manager.getQueueFactoryProvider()
```

```
.createFactory("native", manager.getDatabaseService()));

// Use queue
Queue<Message> queue = factory.createQueue("orders_queue", Message.class);
```

## Integration Tests

```
// Setup creates queue tables dynamically
DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
    .queues(List.of(
        new QueueConfig.Builder()
            .queueName("test_native_queue")
            .maxRetries(3)
            .build()
    ))
    .build();

DatabaseSetupResult result = setupService.createCompleteSetup(request).get();

// Native queue factory available in result
QueueFactory factory = result.getQueueFactories().get("test_native_queue");
```

## Unit Tests

```
// Initialize only native queue schema
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.NATIVE_QUEUE,
    SchemaComponent.DEAD_LETTER_QUEUE
);

// Direct JDBC operations for testing
try (Connection conn = DriverManager.getConnection(jdbcUrl, user, password)) {
    PreparedStatement stmt = conn.prepareStatement(
        "INSERT INTO peegueq.queue_messages (topic, payload) VALUES (?, ?)"
    );
    // ... test operations
}
```

## Key Features:

- Lock-free message processing using `message_processing` table
- Priority-based message ordering (1-10 scale)
- Automatic retry with exponential backoff
- Dead letter queue for failed messages
- Real-time visibility control

## peegeeq-outbox (Transactional Outbox Pattern)

**Purpose:** Transactional message publishing with ACID guarantees (5k+ msg/sec)

### Database Requirements:

- `outbox` table (from Flyway migrations or setup service)
- `outbox_consumer_groups` table (for fanout to multiple consumers)
- `dead_letter_queue` table
- Indexes on status, topic, consumer\_group\_name

### Setup Patterns:

#### Production

```
// Database already initialized via Flyway
PeeGeeQConfiguration config = new PeeGeeQConfiguration();
PeeGeeQManager manager = new PeeGeeQManager(config);

// Get outbox factory
QueueFactory factory = manager.getQueueFactoryProvider()
    .createFactory("outbox", manager.getDatabaseService());

// Outbox pattern: insert into outbox within business transaction
try (Connection conn = dataSource.getConnection()) {
    conn.setAutoCommit(false);

    // Business operation
    stmt.executeUpdate("INSERT INTO orders (id, customer, amount) VALUES (...");

    // Outbox message (same transaction)
    stmt.executeUpdate("INSERT INTO peegeeq.outbox (topic, payload) VALUES
(...");

    conn.commit(); // Both committed atomically
}
```

### Integration Tests

```
// Setup service creates outbox tables
DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
    .queues(List.of()) // No native queues needed
    .eventStores(List.of())
    .build();

DatabaseSetupResult result = setupService.createCompleteSetup(request).get();

// Outbox tables available via PeeGeeQManager
PeeGeeQManager manager = result.getManager();
```

## Unit Tests

```

// Initialize outbox schema
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.OUTBOX,
    SchemaComponent.DEAD_LETTER_QUEUE
);

// Test outbox operations
try (Connection conn = DriverManager.getConnection(jdbcUrl, user, password)) {
    conn.setAutoCommit(false);

    PreparedStatement stmt = conn.prepareStatement(
        "INSERT INTO peegeeq.outbox (topic, payload, status) VALUES (?, ?, ?"
        "'PENDING')"
    );
    stmt.setString(1, "test_topic");
    stmt.setString(2, "{\"data\":\"test\"}");
    stmt.executeUpdate();

    conn.commit();
}

```

## Key Features:

- ACID transaction guarantees with business data
- Consumer group fanout (one message → multiple consumers)
- Status tracking per consumer group
- Automatic retry with configurable backoff
- Dead letter queue for permanent failures

## peegeeq-bitemporal (Bi-temporal Event Store)

**Purpose:** Event sourcing with temporal queries and corrections (3k+ msg/sec)

### Database Requirements:

- `bitemporal_event_log` table (from Flyway migrations or setup service)
- Event store tables created dynamically per event store
- Indexes for temporal queries (`valid_time`, `transaction_time`)
- Triggers for notification streaming
- Views for current state and statistics

### Setup Patterns:

#### Production

```

// Database initialized via Flyway
PeeGeeQConfiguration config = new PeeGeeQConfiguration();
PeeGeeQManager manager = new PeeGeeQManager(config);

// Create event store factory
EventStoreFactory factory = new PeeGeeQEventStoreFactory(manager);

// Create event store for specific entity
EventStore<OrderEvent> orderEvents = factory.createEventStore(
    OrderEvent.class,
    "bitemporal.order_events" // Table name
);

// Append events
BiTemporalEvent<OrderEvent> event = orderEvents.append(
    "order-123",           // event ID
    "OrderCreated",        // event type
    Instant.now(),          // valid time
    orderCreatedPayload    // event data
).get();

```

## Integration Tests

```

// Setup creates event store tables dynamically
DatabaseSetupRequest request = new DatabaseSetupRequest.Builder()
    .eventStores(List.of(
        new EventStoreConfig.Builder()
            .eventStoreName("test_order_events")
            .tableName("test_order_event_log")
            .notificationPrefix("order_event_")
            .build()
    ))
    .build();

DatabaseSetupResult result = setupService.createCompleteSetup(request).get();

// Event stores available in result
EventStore<?> eventStore = result.getEventStores().get("test_order_events");

```

## Unit Tests

```

// Initialize bitemporal schema
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.BITEMPORAL
);

// Test direct SQL operations

```

```

try (Connection conn = DriverManager.getConnection(jdbcUrl, user, password)) {
    PreparedStatement stmt = conn.prepareStatement(
        "INSERT INTO bitemporal.bitemporal_event_log " +
        "(event_id, event_type, valid_time, transaction_time, payload) " +
        "VALUES (?, ?, ?, ?, ?::jsonb)"
    );
    stmt.setString(1, "evt-123");
    stmt.setString(2, "TestEvent");
    stmt.setTimestamp(3, Timestamp.from(Instant.now()));
    stmt.setTimestamp(4, Timestamp.from(Instant.now()));
    stmt.setString(5, "{\"data\":\"test\"}");
    stmt.executeUpdate();
}

```

**Key Features:**

- Bi-temporal tracking (valid\_time + transaction\_time)
- Event corrections with audit trail
- Point-in-time queries (as of any timestamp)
- Aggregate versioning and history
- JSONB payload for flexible event data
- Notification triggers for event streaming

**Temporal Queries:**

```

-- Current state (as of now)
SELECT * FROM bitemporal_current_state WHERE event_id = 'order-123';

-- Historical state (as of specific time)
SELECT * FROM bitemporal_event_log
WHERE event_id = 'order-123'
    AND transaction_time <= '2025-11-01'
ORDER BY transaction_time DESC
LIMIT 1;

-- All corrections for an event
SELECT * FROM bitemporal_event_log
WHERE event_id = 'order-123'
ORDER BY version;

```

**peegeeq-migrations (Flyway Migration Module)**

**Purpose:** Production database initialization and version control

**Setup Method:** Flyway migrations only (this IS the migration module)

**Key Files:**

- [src/main/resources/db/migration/V001\\_\\_Create\\_Base\\_Tables.sql](#) (576 lines)
  - Complete schema definition

- All tables, indexes, functions, triggers
- Idempotent (safe to re-run)

## Usage:

### Standalone JAR

```
# Build  
mvn clean package -DskipTests  
  
# Run migrations  
java -jar target/peegeeq-migrations.jar migrate  
  
# Check status  
java -jar target/peegeeq-migrations.jar info  
  
# Validate  
java -jar target/peegeeq-migrations.jar validate
```

### Maven Plugin

```
cd peegeeq-migrations  
  
# Run migrations  
mvn flyway:migrate  
  
# Check status  
mvn flyway:info  
  
# Clean database (dev only!)  
mvn flyway:clean  
  
# Validate migrations  
mvn flyway:validate
```

### Docker

```
# Build image  
docker build -t peegeeq-migrations .  
  
# Run migrations  
docker run --rm \  
-e DB_JDBC_URL=jdbc:postgresql://host:5432/db \  
-e DB_USER=user \  
-e DB_PASSWORD=pass \  
peegeeq-migrations migrate
```

## Migration Script Structure:

```
-- V001__Create_Base_Tables.sql

-- Header (metadata)
-- executeInTransaction=false

-- Schema version tracking
CREATE TABLE IF NOT EXISTS schema_version (...);

-- Core tables
CREATE TABLE IF NOT EXISTS outbox (...);
CREATE TABLE IF NOT EXISTS queue_messages (...);
CREATE TABLE IF NOT EXISTS bitemporal_event_log (...);
-- ... all tables

-- Indexes
CREATE INDEX idx_outbox_status_created ON outbox(status, created_at);
CREATE INDEX idx_queue_messages_topic_visible ON queue_messages(topic,
visible_at);
-- ... 30+ indexes

-- Views
CREATE OR REPLACE VIEW bitemporal_current_state AS ...;
-- ... 3 views

-- Functions and Triggers
-- (defined in migration script)
```

## peegeeq-test-support (Test Utilities)

**Purpose:** Shared test utilities for all modules

### Key Classes:

- **PeeGeeQTestSchemaInitializer** (747 lines)
  - Component-based schema initialization
  - Fast JDBC-based setup for unit tests
  - Granular schema component control

### Usage Patterns:

#### Full Schema

```
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.ALL
);
```

## Selective Components

```
// Only what you need for specific test
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.OUTBOX,
    SchemaComponent.NATIVE_QUEUE,
    SchemaComponent.DEAD_LETTER_QUEUE
);
```

## Queue-Related Only

```
// Convenience enum for all queue components
PeeGeeQTestSchemaInitializer.initializeSchema(
    postgres,
    SchemaComponent.QUEUE_ALL // Expands to OUTBOX + NATIVE + DLQ
);
```

## Cleanup Between Tests

```
@AfterEach
void cleanup() {
    // Remove test data but keep schema
    PeeGeeQTestSchemaInitializer.cleanupTestData(
        postgres,
        SchemaComponent.OUTBOX,
        SchemaComponent.NATIVE_QUEUE
    );
}
```

## Advantages Over PeeGeeQDatabaseSetupService:

- **Faster:** Direct JDBC, no Vert.x startup
- **Simpler:** No database creation, uses TestContainer database
- **Granular:** Initialize only what you need
- **Synchronous:** No CompletableFuture complexity
- **Ideal for unit tests:** Testing specific components

## Database Schema Components

### Complete Schema Overview

**Note:** PeeGeeQ uses **two separate PostgreSQL schemas** by default (`peegueq` and `bitemporal`), but both can be customized to any schema name(s) you prefer. You can even use the same schema for both (e.g.,

myapp for everything) or integrate into existing schemas.

## PeeGeeQ Database Schema

### Schemas (2) - Fully Customizable

- peegueq (main schema) - configurable
- bitemporal (event store schema) - configurable

### Core Tables (7)

- schema\_version
- queue\_messages (native queue)
- outbox (transactional outbox)
- outbox\_consumer\_groups (fanout)
- message\_processing (lock-free)
- dead\_letter\_queue (errors)
- bitemporal\_event\_log (event sourcing)

### Template Tables (2)

- peegueq.queue\_template
- bitemporal.event\_store\_template

### Monitoring Tables (2)

- queue\_metrics
- connection\_pool\_metrics

### Consumer Group Tables (5)

- peegueq.consumer\_group
- peegueq.consumer\_instance
- peegueq.consumer\_subscription
- peegueq.consumer\_offset
- peegueq.consumer\_heartbeat

### Indexes (30+)

- Performance indexes on all tables
- Composite indexes for queries
- GIN indexes for JSONB columns
- Unique indexes for constraints

### Functions & Triggers (4)

- Timestamp update triggers
- Notification triggers

#### Views (3)

- bitemporal\_current\_state
- bitemporal\_latest\_events
- bitemporal\_event\_stats

## Table Relationships

```

queue_messages    --> message_processing (1:N)
outbox           --> dead_letter_queue (1:N)
outbox           -----> outbox_consumer_groups (1:N)

queue_template <-- queue tables (inheritance)
event_store_template <-- event store tables (inheritance)

bitemporal_event_log --> bitemporal_current_state (view)
                      |--> bitemporal_latest_events (view)
                      |--> bitemporal_event_stats (view)

```

## Schema Component Dependencies

```

Extensions (uuid-ossp, pg_stat_statements)
  ↓
Schemas (peegeeq, bitemporal)
  ↓
Template Tables (queue_template, event_store_template)
  ↓
Core Tables (queue_messages, outbox, bitemporal_event_log)
  ↓
Supporting Tables (dead_letter_queue, message_processing)
  ↓
Consumer Group Tables (5 tables)
  ↓
Indexes (30+)
  ↓
Functions & Triggers (4)
  ↓
Views (3)

```

# Setup Methods Comparison

## Feature Matrix

Feature	Flyway Migrations	PeeGeeQDatabaseSetupService	PeeGeeQTestSchemaInitializer
Environment	Production, Dev	Integration Tests	Unit Tests
Speed	Medium (1-2s)	Slow (5-10s)	Fast (<1s)
Complexity	Low	High	Low
Flexibility	Low	High	Medium
Version Control	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No
Rollback	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No
Parameterization	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Dynamic DB Creation	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Async/Reactive	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No
Component Selection	<input type="checkbox"/> No	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes
Cleanup	Manual	<input checked="" type="checkbox"/> Automatic	<input checked="" type="checkbox"/> Automatic

## When to Use Each Method

### Use Flyway Migrations When:

- Setting up production databases
- Need version control and audit trail
- Schema changes need to be tracked
- Multiple environments need same schema
- Rollback capability is required
- CI/CD pipeline deployment

### Use PeeGeeQDatabaseSetupService When:

- Integration tests need isolated databases
- Dynamic database creation per test
- Parameterized table names required
- Testing full PeeGeeQManager lifecycle
- Need QueueFactory and EventStore instances
- Async/reactive patterns being tested

## Use PeeGeeQTestSchemaInitializer When:

- Unit tests need minimal schema setup
  - Speed is critical (fast test execution)
  - Only specific components needed
  - Testing individual components in isolation
  - No PeeGeeQManager instance needed
  - Simple JDBC operations being tested
- 

## Troubleshooting

### Common Issues

#### Issue 1: "Table does not exist" Errors

##### Symptoms:

```
FATAL: dead_letter_queue table does not exist - schema not initialized properly
FATAL: queue_messages table does not exist - schema not initialized properly
```

**Cause:** Database not initialized before application startup

##### Solution:

```
# Production: Run migrations first
cd peegeeq-migrations
mvn flyway:migrate

# Development: Use docker-compose
docker-compose -f docker-compose.dev.yml up

# Integration Tests: Use PeeGeeQDatabaseSetupService
setupService.createCompleteSetup(request).get();

# Unit Tests: Use PeeGeeQTestSchemaInitializer
PeeGeeQTestSchemaInitializer.initializeSchema(postgres, SchemaComponent.ALL);
```

#### Issue 2: "Templates not found" in Tests

##### Symptoms:

```
Templates not found in database testdb: queue_template=false,
event_store_template=false
```

**Cause:** Extensions not created (missing uuid-ossp or pg\_stat\_statements)

**Solution:**

```
// Ensure extensions are available in PostgreSQL image
@Container
private static PostgreSQLContainer<?> postgres = new PostgreSQLContainer<>(
    "postgres:15.13-alpine3.20")
    .withCommand("postgres -c shared_preload_libraries=pg_stat_statements");

// OR use base template which handles extension failures gracefully
```

### Issue 3: Multi-Statement SQL Not Executing

**Symptoms:**

- Only first SQL statement in file executes
- Missing indexes, triggers, or tables

**Cause:** Vert.x PostgreSQL client limitation (only executes first statement)

**Solution:** SQL templates are now organized as directories with numbered files

```
base/
└── .manifest           # Lists files in order
└── 01a-extension.sql   # One statement per file
└── 01b-schema.sql
└── ...
```

**See Also:** [docs/devtest/VERTX\\_MULTISTATEMENT\\_SQL\\_BUG\\_ANALYSIS.md](#)

### Issue 4: Parallel Test Failures

**Symptoms:**

```
Database "test_db_123" is being accessed by other users
Test 4's manager is connecting to Test 5's database
```

**Cause:** Shared database name or non-unique setup IDs

**Solution:**

```
// Use unique database name per test class
String dbName = "test_" + getClass().getSimpleName() + "_" + UUID.randomUUID();

// Use unique setup ID per test
```

```
String setupId = "test_" + Thread.currentThread().getId() + "_" +
System.nanoTime();

// Use programmatic configuration (not System.setProperty)
PeeGeeQConfiguration config = new PeeGeeQConfiguration(
    setupId,
    host,
    port,
    dbName,
    username,
    password,
    schema
);
```

## Issue 5: Flyway Migration Conflicts

### Symptoms:

```
Validate failed: Migrations have failed validation
Detected resolved migration not applied to database: 1
```

**Cause:** Database state doesn't match migration history

### Solution:

```
# Check migration status
mvn flyway:info

# Repair migration history
mvn flyway:repair

# Clean and re-migrate (DEV ONLY!)
mvn flyway:clean flyway:migrate

# Production: Never use clean, investigate and fix manually
```

## Issue 6: Permission Denied on CREATE EXTENSION

### Symptoms:

```
ERROR: permission denied to create extension "uuid-ossp"
ERROR: must be superuser to create extension "pg_stat_statements"
```

**Cause:** Database user lacks CREATE EXTENSION permission

**Solution:**

```
-- Grant extension creation permission
ALTER USER peegeeq_user WITH SUPERUSER;

-- OR create extensions as superuser first
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";
CREATE EXTENSION IF NOT EXISTS "pg_stat_statements";

-- Then run migrations as regular user
```

**Alternative:** Production code has fallback to minimal schema without extensions

**Issue 7: Test Cleanup Failures****Symptoms:**

```
Failed to drop test database: test_db_123 - database is being accessed by other
users
```

**Cause:** Open connections still active

**Solution:**

```
@AfterAll
void cleanup() {
    // Stop PeeGeeQManager first (closes all connections)
    if (setupResult != null && setupResult.getManager() != null) {
        setupResult.getManager().stop();
    }

    // Then destroy setup
    if (setupService != null && setupResult != null) {
        setupService.destroySetup(setupResult.getSetupId()).get();
    }
}
```

**Debug Logging**

Enable debug logging to troubleshoot setup issues:

```
# logback.xml or application.properties
logging.level.dev.mars.peegeeq.db.setup=DEBUG
logging.level.org.flywaydb=DEBUG
```

## Key log messages:

```

🔧 Applying base template: base
🔧 ✅ Base template SQL executed
🔧 🔎 Verifying templates exist in bitemporal schema...
✅ FINAL: All schema templates applied successfully

```

## Summary

### Quick Reference

Task	Method	Command/Code
<b>Setup</b>		
<b>Production</b>	Flyway Migrations	<code>java -jar peegeeq-migrations.jar migrate</code>
<b>DB</b>		
<b>Setup Dev</b>		
<b>DB</b>	Docker Compose	<code>docker-compose -f docker-compose.dev.yml up</code>
<b>Setup</b>		
<b>Integration</b>	PeeGeeQDatabaseSetupService	<code>setupService.createCompleteSetup(request).get()</code>
<b>Test</b>		
<b>Setup Unit</b>		
<b>Test</b>	PeeGeeQTestSchemaInitializer	<code>initializeSchema(postgres, SchemaComponent.ALL)</code>
<b>Set Custom</b>		
<b>Schema</b> <b>(Flyway)</b>	Configuration	<code>flyway.schemas=myapp_queue,myapp_events</code>
<b>Set Custom</b>		
<b>Schema</b> <b>(App)</b>	Environment Variable	<code>export PEEGEEQ_DB_SCHEMA="myapp_queue"</code>
<b>Verify</b>		
<b>Schema</b>	Flyway Info	<code>mvn flyway:info</code>
<b>Reset Dev</b>		
<b>DB</b>	Flyway Clean	<code>mvn flyway:clean flyway:migrate</code>
<b>Check</b>		
<b>Health</b>	PeeGeeQManager	<code>manager.getHealthService().checkHealth()</code>

### Key Takeaways

- Production = Flyway Migrations:** Always use `peegeeq-migrations` module
- Tests = Programmatic Setup:** Use setup service or test initializer
- Custom Schemas Supported:** Deploy to any PostgreSQL schema, not just `public`

4. **Multi-Tenant Ready:** Use schema-per-tenant for isolation and scalability
  5. **One Statement Per File:** Vert.x limitation requires single-statement SQL files
  6. **Unique Database Names:** Critical for parallel test execution
  7. **Cleanup After Tests:** Always stop manager before destroying setup
  8. **Component-Based Testing:** Only initialize what you need in unit tests
  9. **Template-Based Integration:** Use templates for dynamic test databases
  10. **Fully Qualified Names:** Always use schema.table in queries for clarity
- 

## Related Documentation

- [\*\*VERTX\\_MULTISTATEMENT\\_SQL\\_BUG\\_ANALYSIS.md\*\*](#) - Deep dive into Vert.x SQL execution limitation
  - [\*\*PEEGEEQ.Migrations\\_DEPLOYMENT\\_GUIDE.md\*\*](#) - Production deployment patterns
  - [\*\*PEEGEEQ.Migrations\\_README.md\*\*](#) - Migration module overview
  - [\*\*TESTING-GUIDE.md\*\*](#) - Comprehensive testing guide
  - [\*\*PEEGEEQ\\_COMPLETE\\_GUIDE.md\*\*](#) - Full system architecture
- 

© Mark Andrew Ray-Smith Cityline Ltd 2025